COST EFFECTIVE KUBERNETES

Myroslav Rys — Softserve
Ryan Richard — Google

softserve

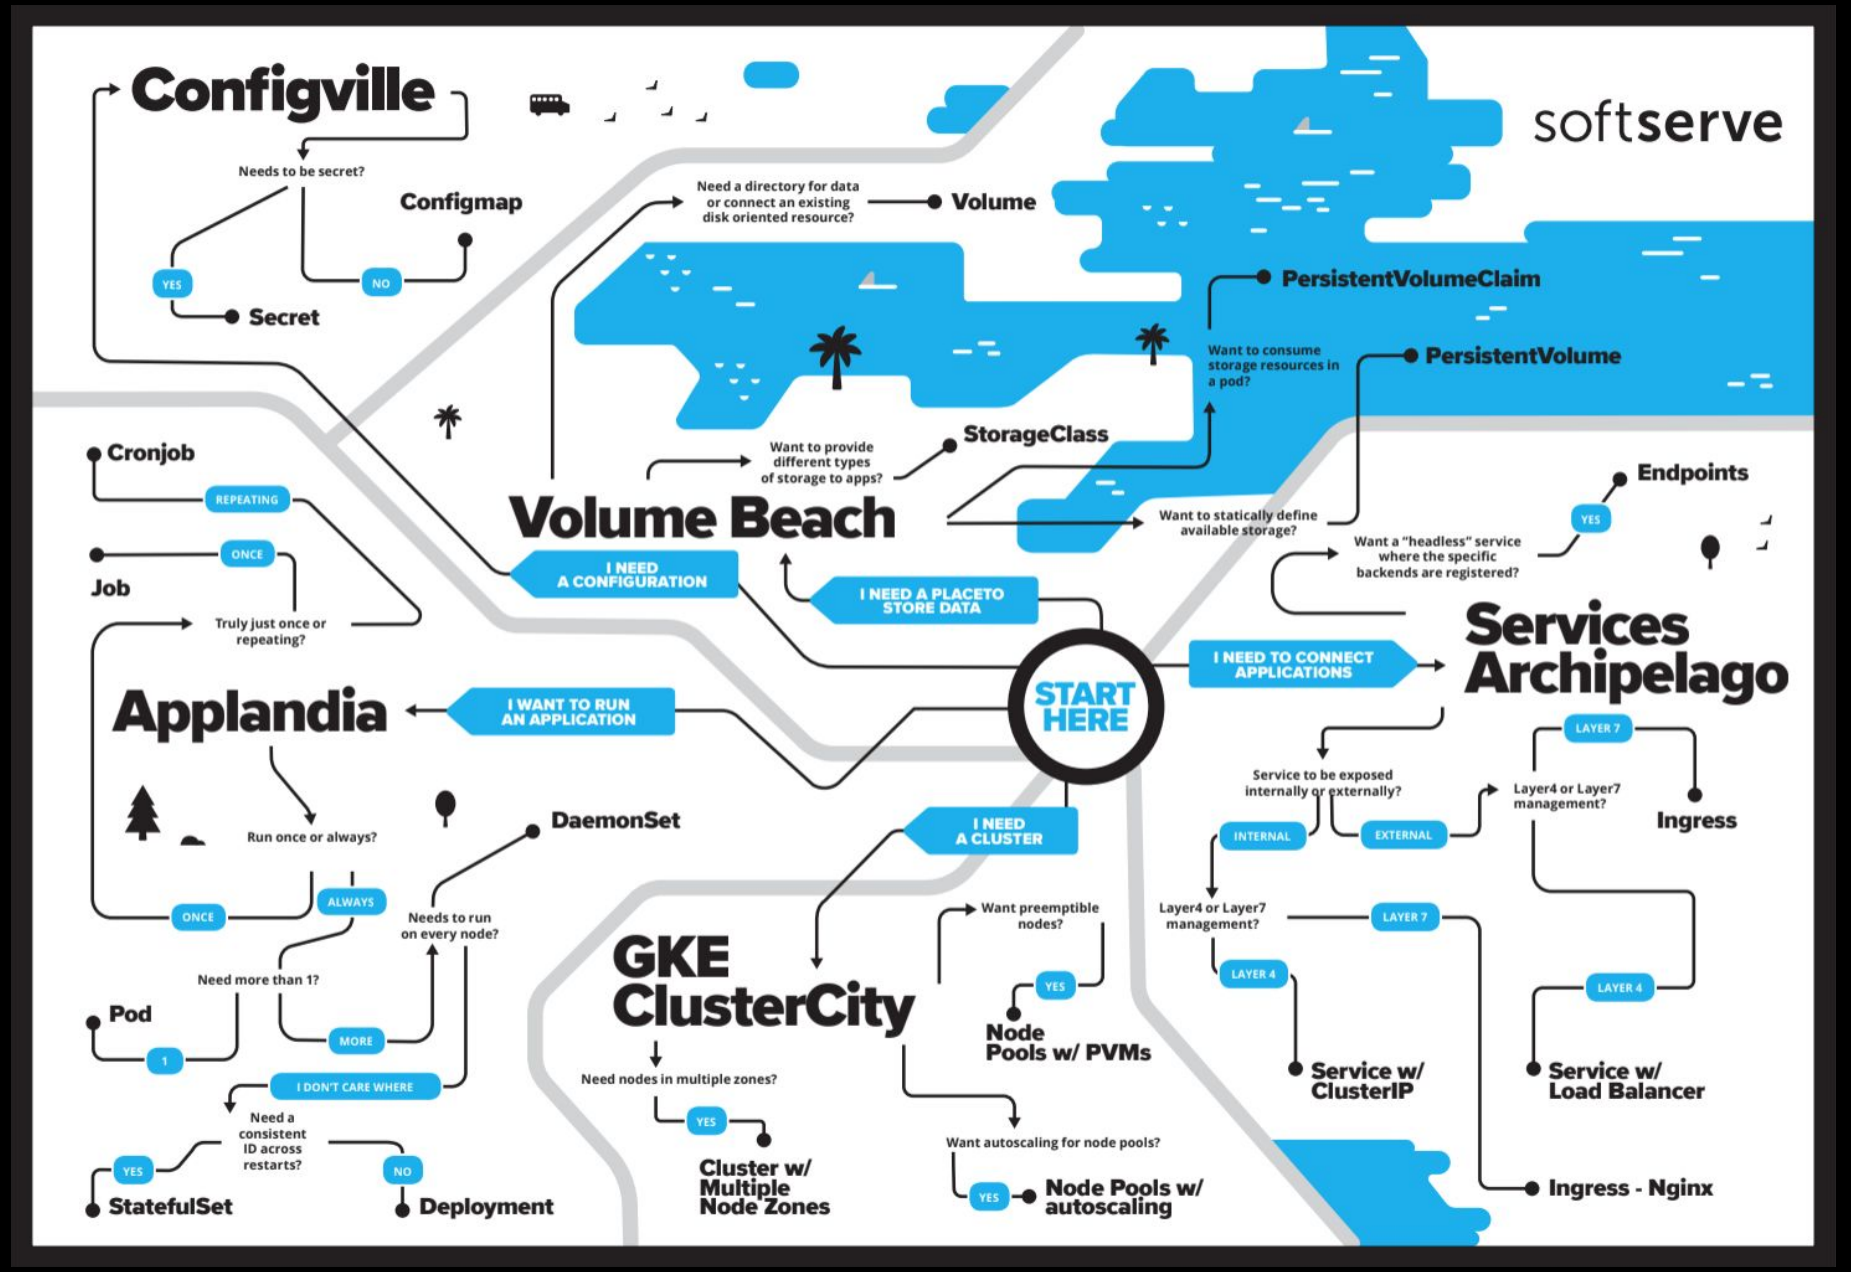
Configville
softserve
Needs to be secret?
Configmap
YES
NO
Secret
Need a directory for data or connect an existing disk oriented resource?
Volume
PersistentVolumeClaim
Want to consume storage resources in a pod?
PersistentVolume
Want to provide different types of storage to apps?
StorageClass
Cronjob
REPEATING
Job
ONCE
Truly just once or repeating?
Volume Beach
Want to statically define available storage?
Endpoints
YES
Want a "headless" service where the specific backends are registered?
I NEED A CONFIGURATION
I NEED A PLACE TO STORE DATA
I NEED TO CONNECT APPLICATIONS
Services Archipelago
START HERE
Applandia
I WANT TO RUN AN APPLICATION
Run once or always?
DaemonSet
ONCE
ALWAYS
Needs to run on every node?
Need more than 1?
Pod
1
MORE
I DON'T CARE WHERE
Need a consistent ID across restarts?
YES
NO
StatefulSet
Deployment
I NEED A CLUSTER
GKE ClusterCity
Want preemptible nodes?
YES
Node Pools w/ PVMs
Need nodes in multiple zones?
YES
Cluster w/ Multiple Node Zones
Want autoscaling for node pools?
YES
Node Pools w/ autoscaling
Service to be exposed internally or externally?
INTERNAL
EXTERNAL
Layer4 or Layer7 management?
LAYER 7
Ingress
Layer4 or Layer7 management?
LAYER 7
LAYER 4
Service w/ ClusterIP
Ingress - Nginx
LAYER 4
Service w/ Load Balancer

softserve

KUBERNETES ARCHITECTURE FOUNDATION

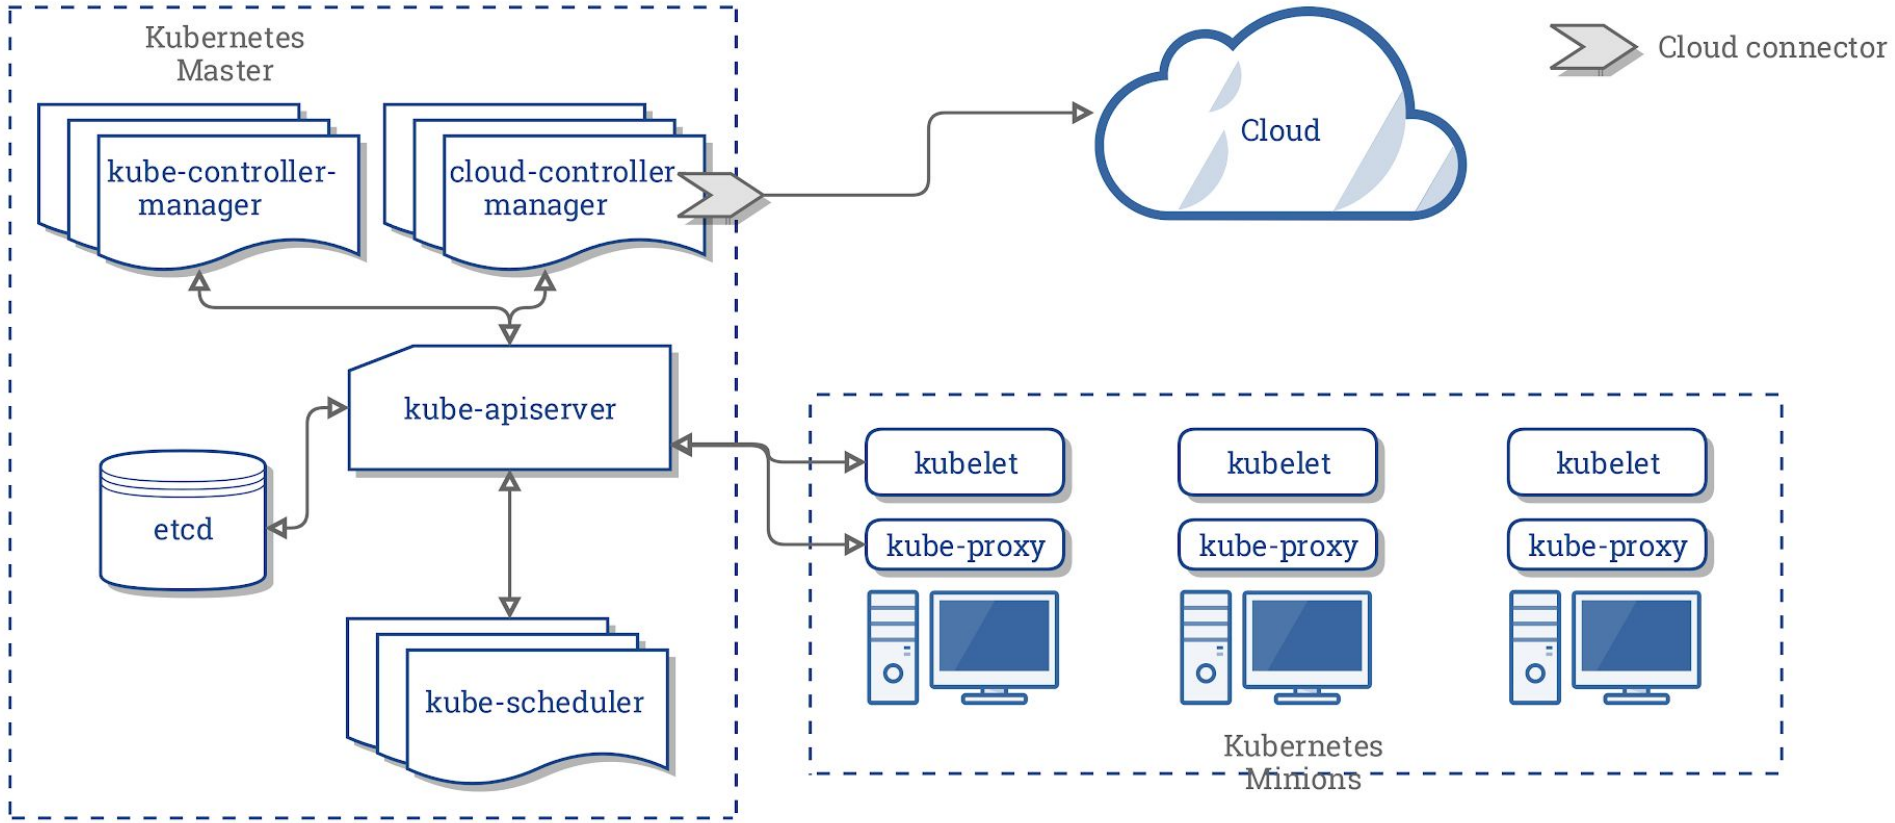

GKE — SINGLE ZONE / NODE POOL

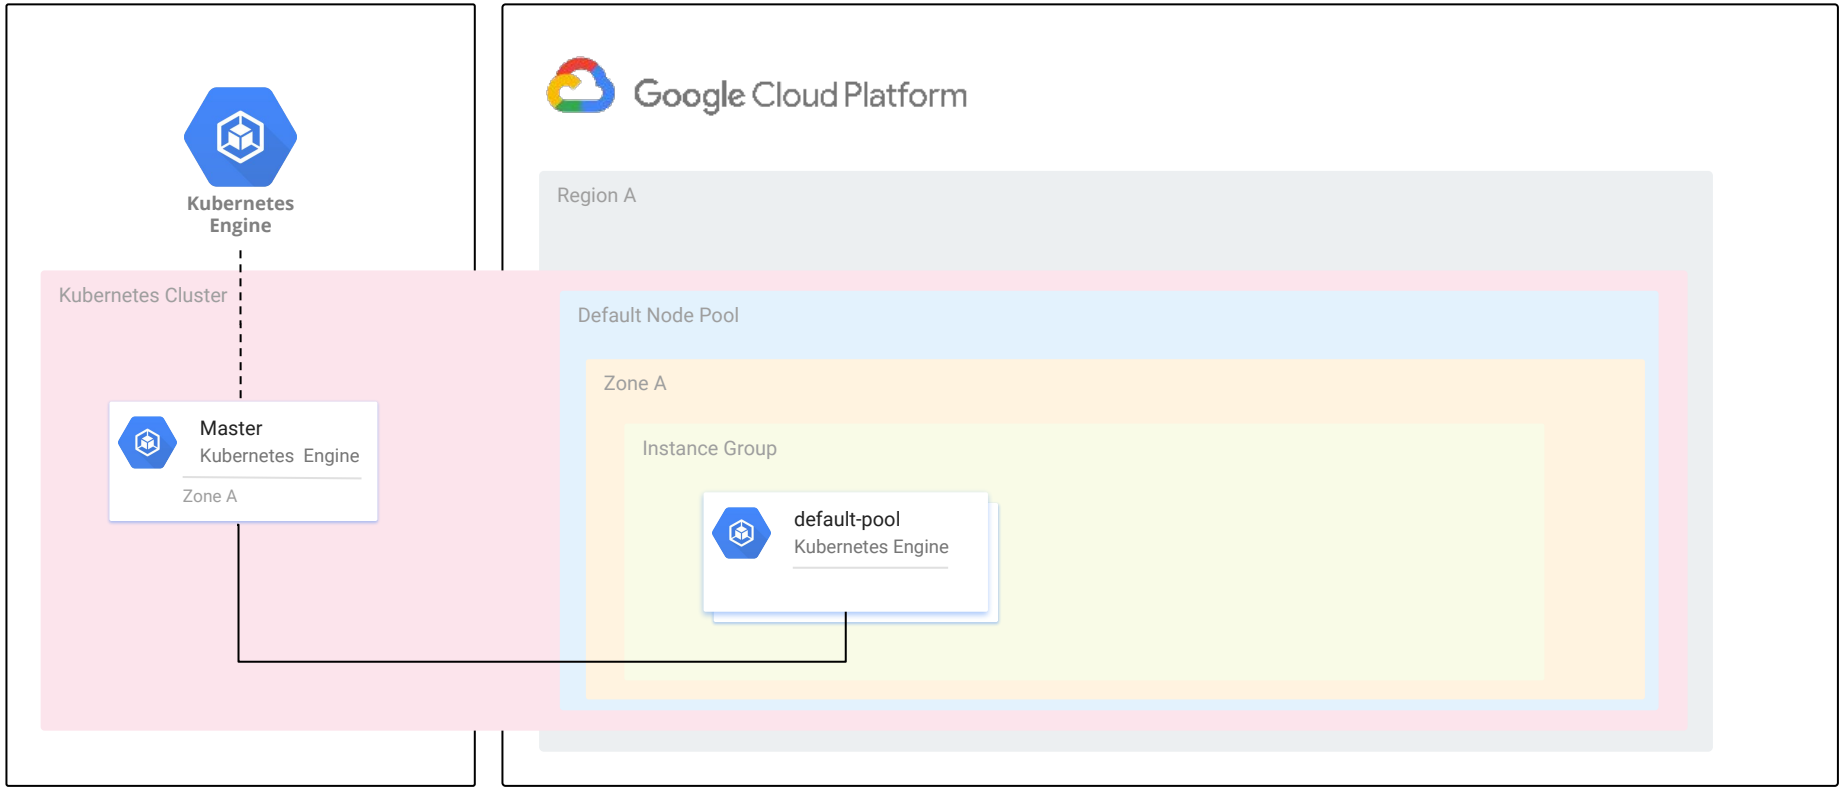

softserve

GKE — SINGLE NODE POOL, MULTI-ZONE

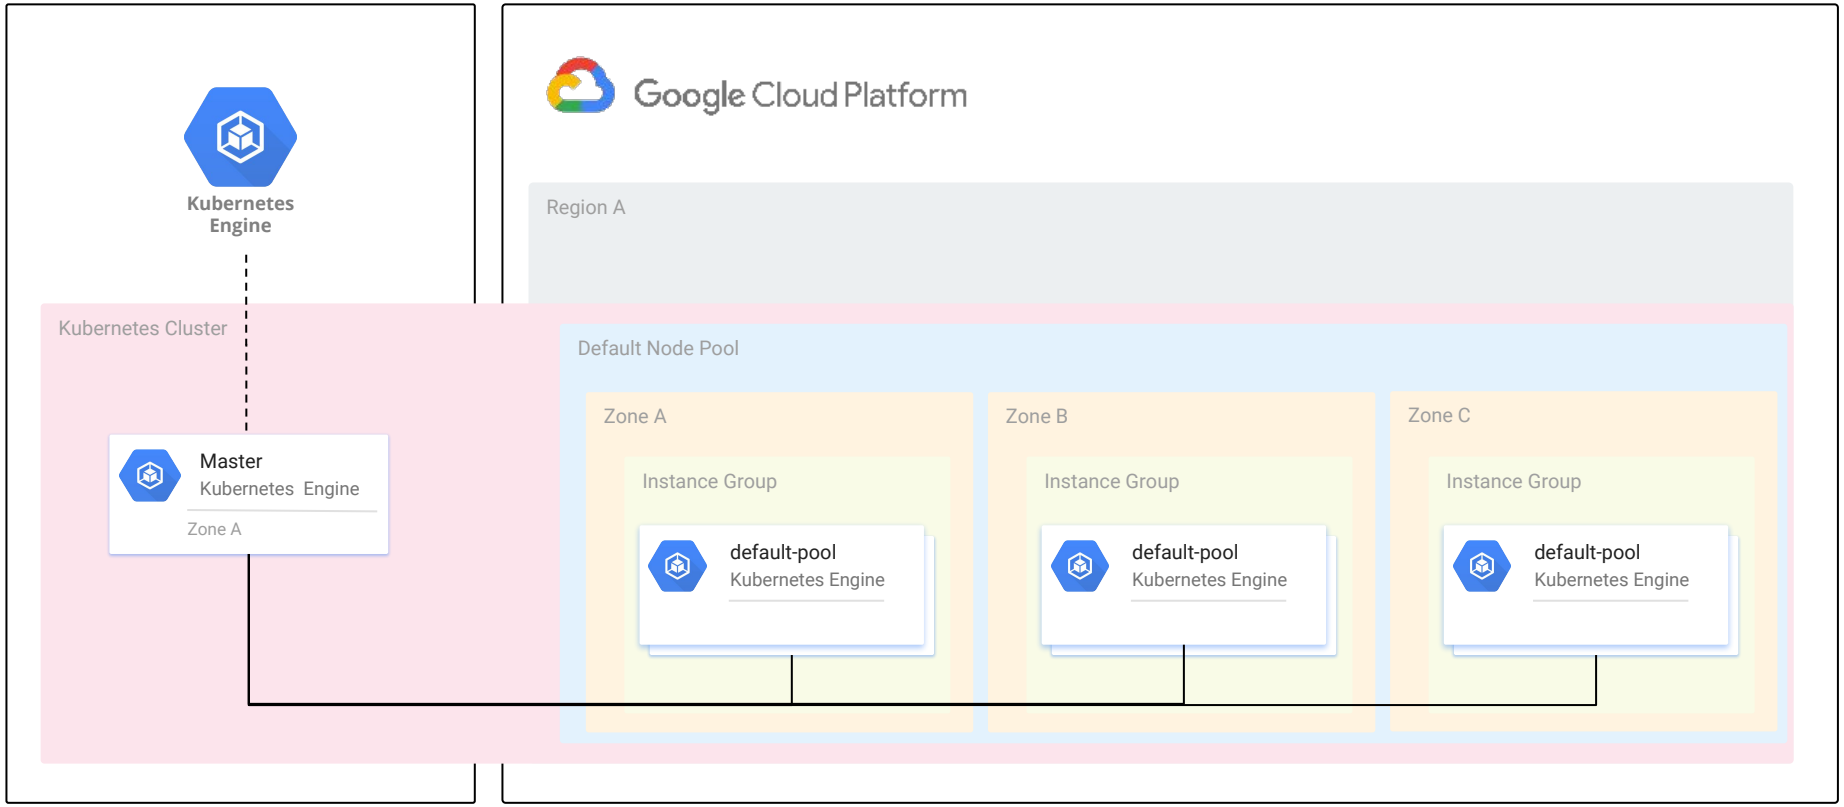

softserve

CAPACITY vs COST

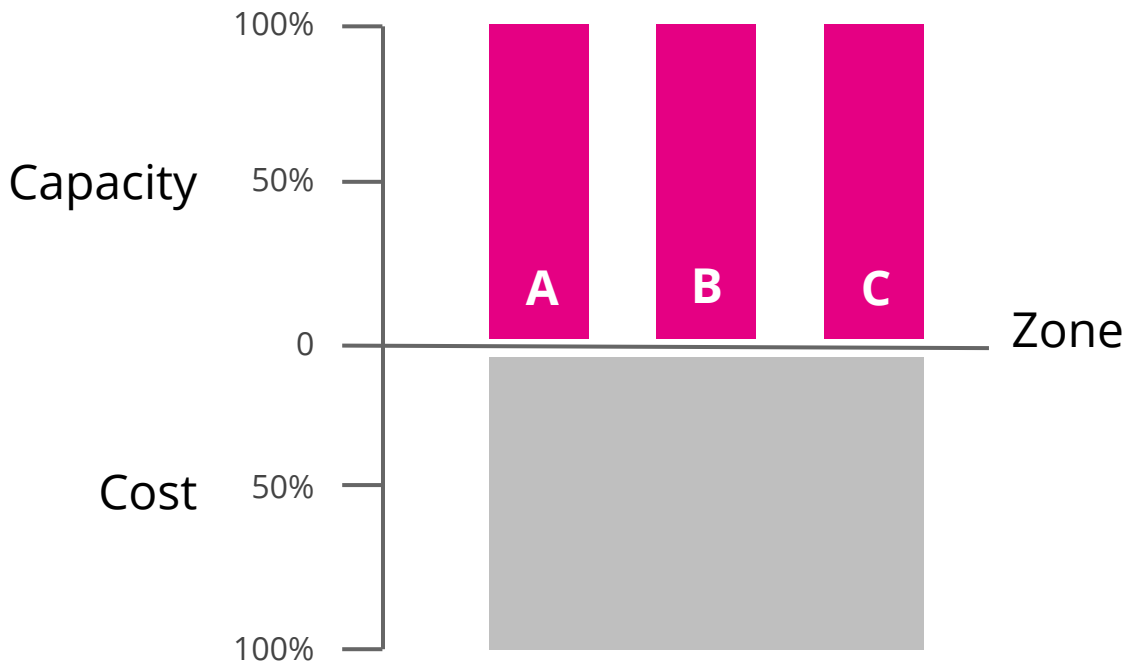

100% of spend

100% of the time

regular instances

softserve

GKE CLUSTER AUTOSCALER

Allows for Autoscaling of the compute nodes which make up the cluster.

gcloud beta container clusters create meetup \

**--enable-autoscaling **

**--min-nodes 2 **

--max-nodes 6

...

softserve

CLUSTER AUTOSCALER VS CUSTOM CONTROLLER OR MANUAL MANAGEMENT

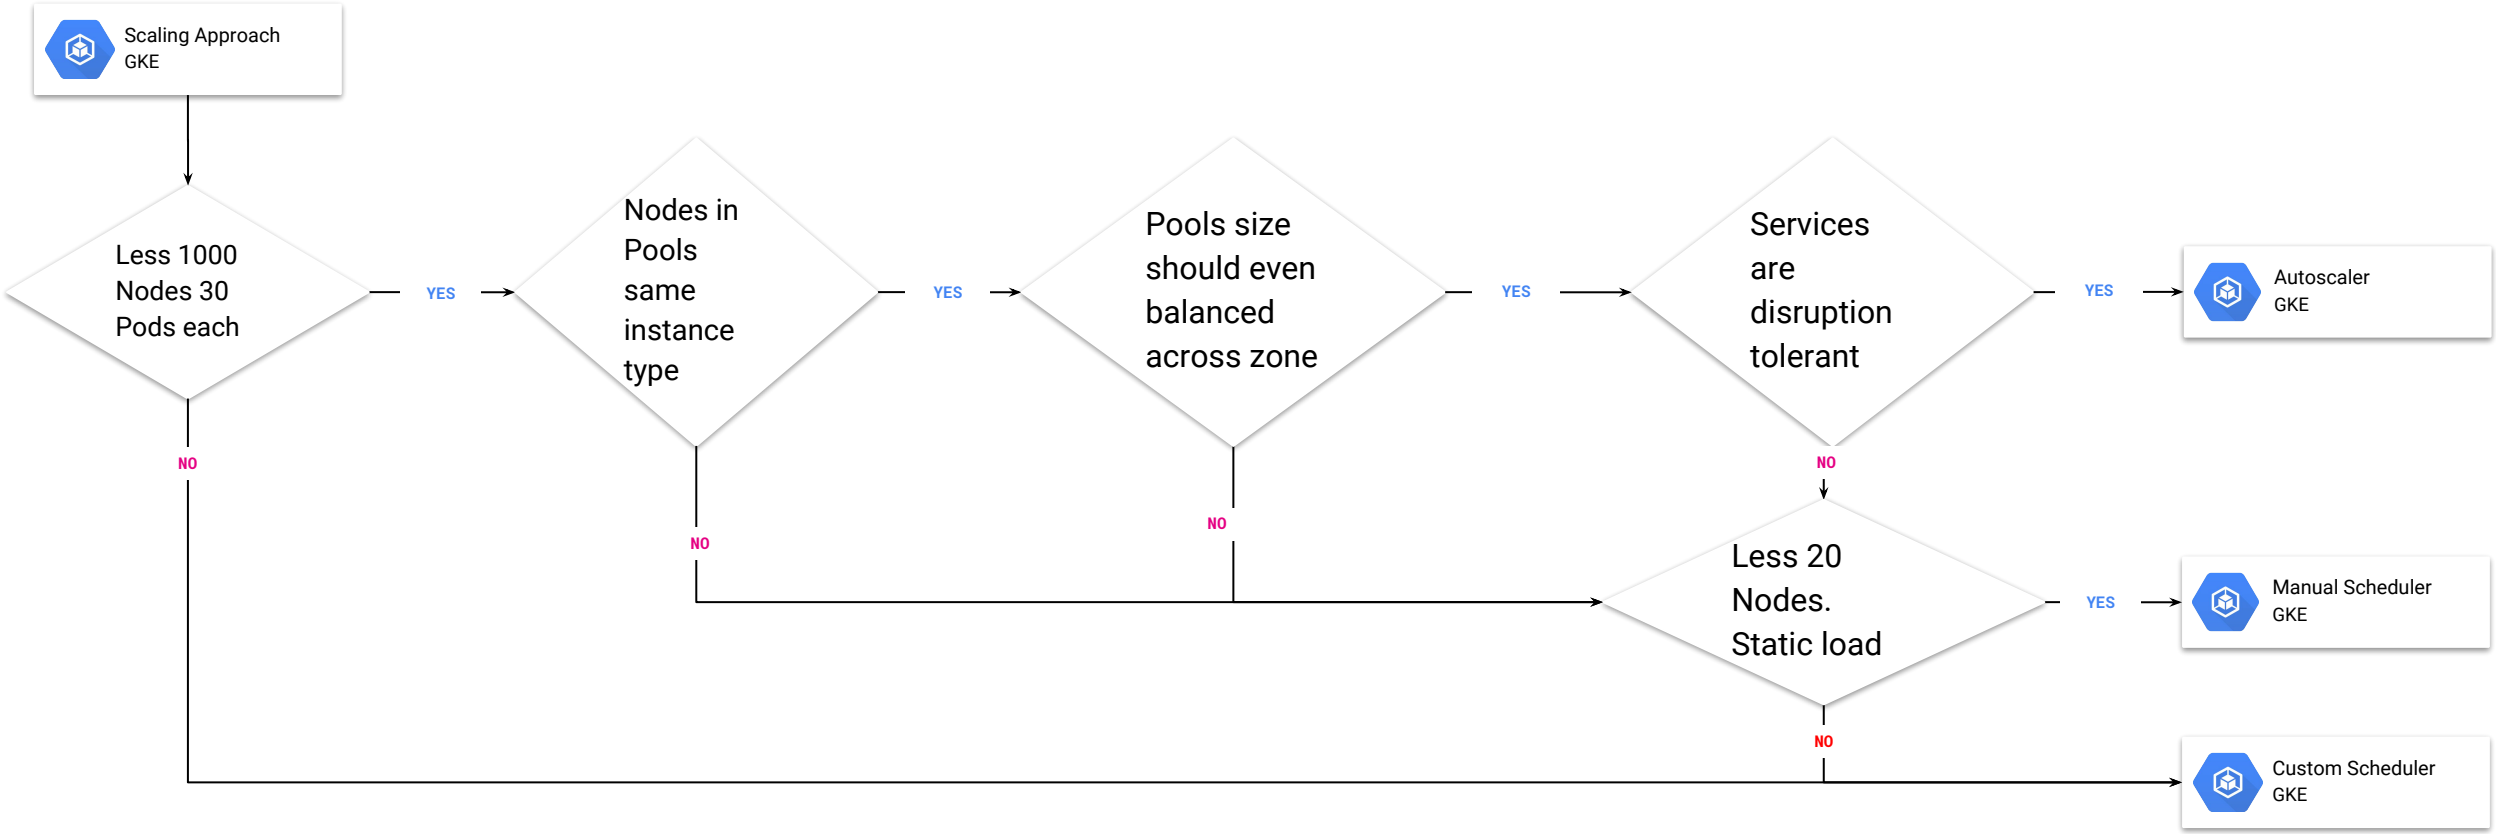

GKE CLUSTER AUTOSCALER LIMITATIONS

- Cluster autoscaler works based on **Pod resource requests**
- Cluster autoscaler does **not track labels manually added** after initial cluster or node pool creation
- Cluster autoscaler considers the relative cost of each instance type in the node pool and attempts to expand the **least expensive** possible node pool
- Cluster with multiple node pools with the same instance type, cluster autoscaler will attempt to keep those node pools' sizes **balanced**
- Maximum period for **graceful termination** for a Pod up to 10 minutes

softserve

CAPACITY VS COST (autoscaling)

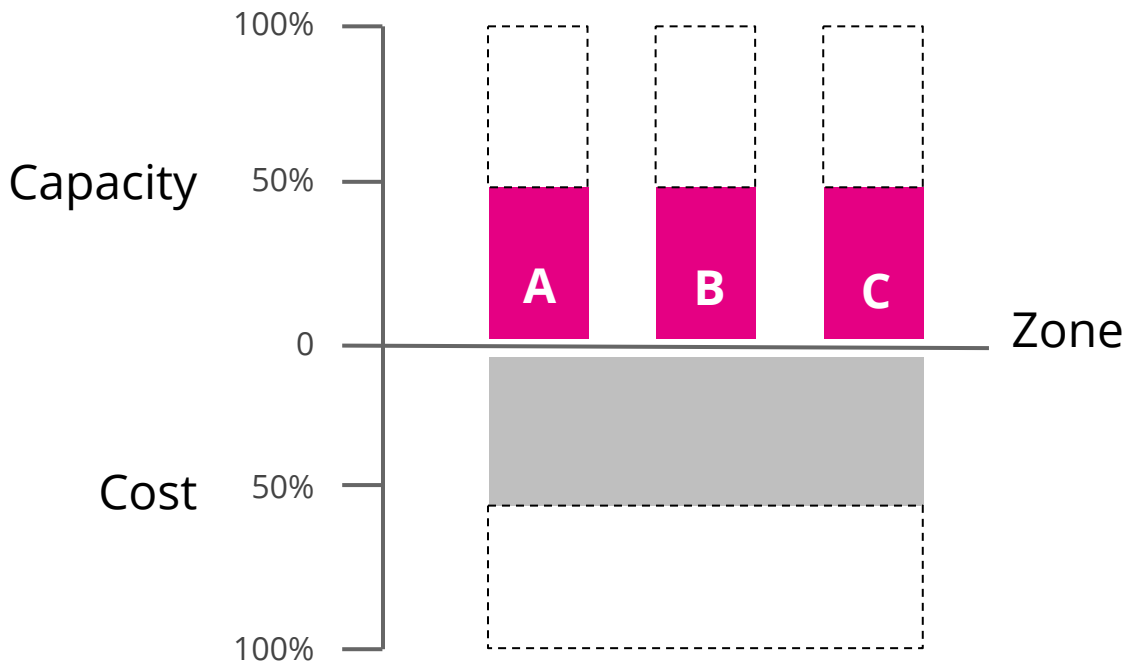

BEST CASE:

50% of spend,
100% of the time

WORST CASE:

100% of spend,
100% of the time

regular instances

softserve

PREEMPTIBLE VMS

A preemptible VM (PVM) is an instance that you can create and run at a much lower price than normal instances. However, Compute Engine might terminate (preempt) these instances if it requires access to those resources for other tasks.

~80% Discount!

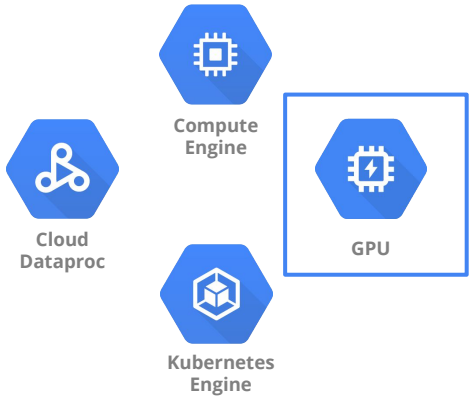

24HR
Life (max)

softserve

PREEMPTIBLE VMS STATS

580,000

20,000

10% - 15%

softserve

PREEMPTIBLE VMS STATS

580,000
cores for 1 HPC workload

$20,000
over a weekend for HPC workload

10% - 15%
Average Preemption Rate*

softserve

EXTEND GKE CLUSTER WITH PVM NODE POOL DEMO

softserve

RECOMMENDED WORKLOAD FOR PVM NODE POOL

```
gcloud container node-pools create pvm-pool \
  --cluster $CLUSTER_NAME \
  --zone $CLUSTER_ZONE \
  --scopes cloud-platform \
  --enable-autoupgrade \
  --preemptible \
  --num-nodes 1 --machine-type g1-small \
  --enable-autoscaling --min-nodes=1 --max-nodes=6
```

softserve

GKE — MULTI-ZONE / MULTIPLE NODE POOLS

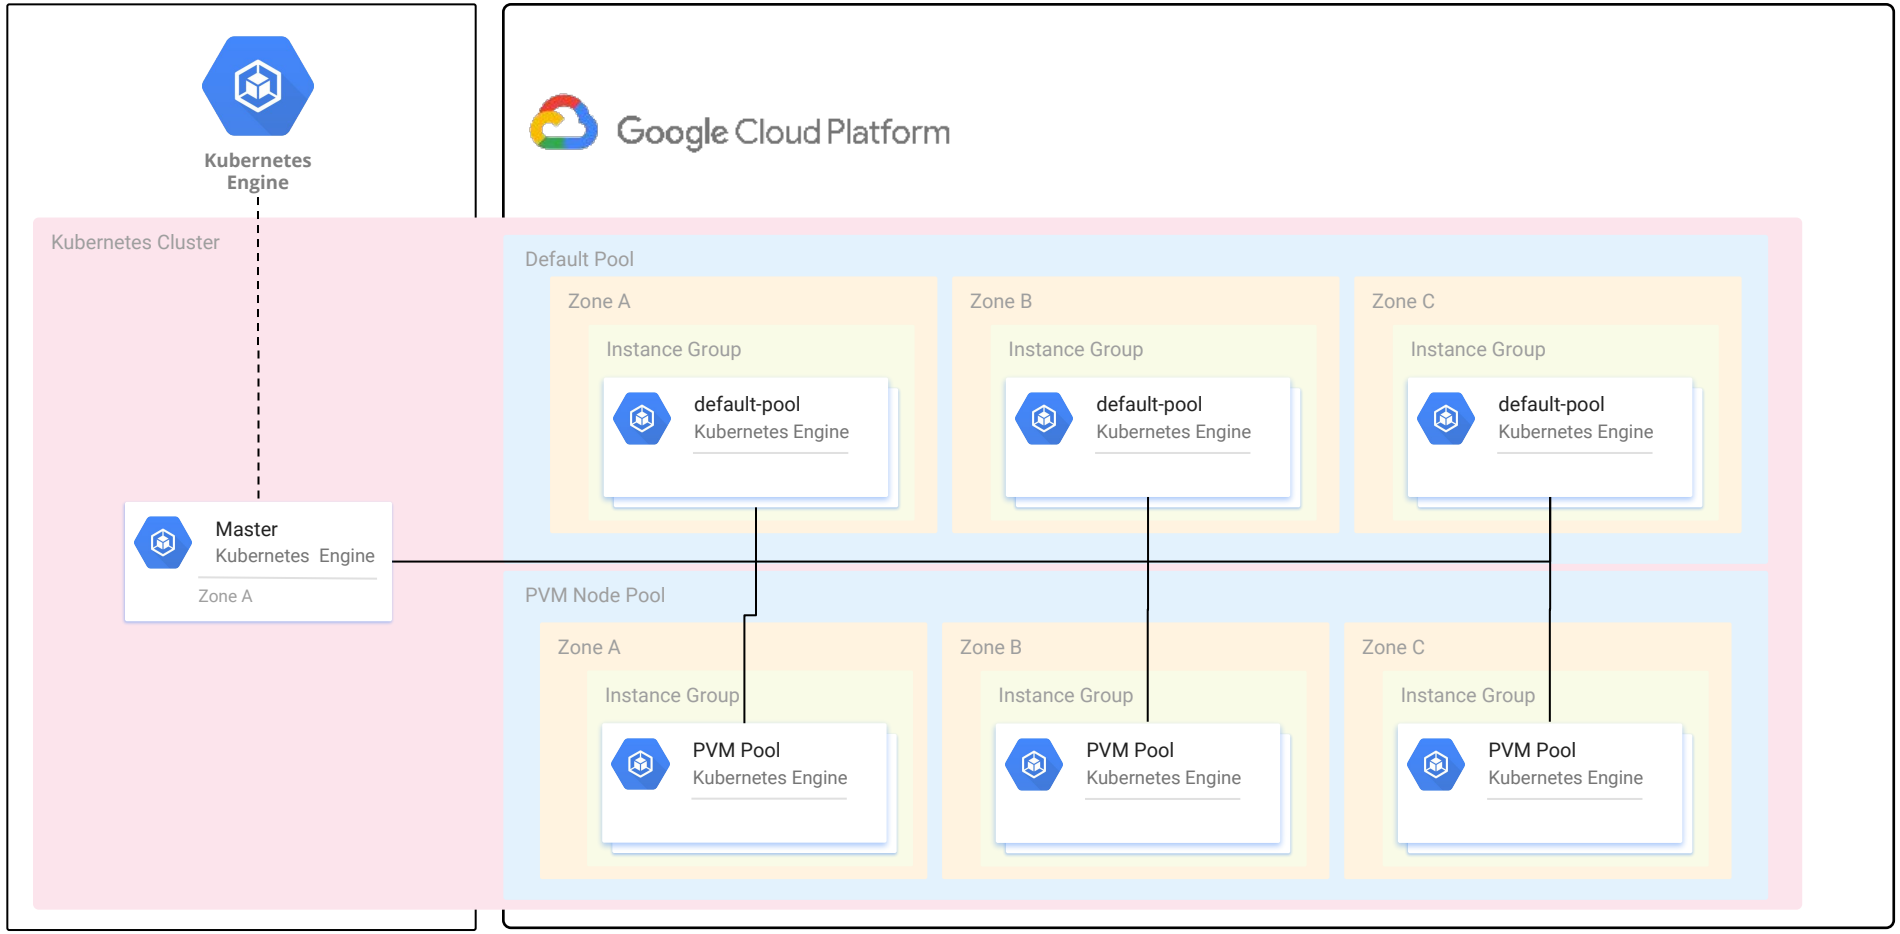

GKE — PREEMPTIBLE POOL DECISION MAP

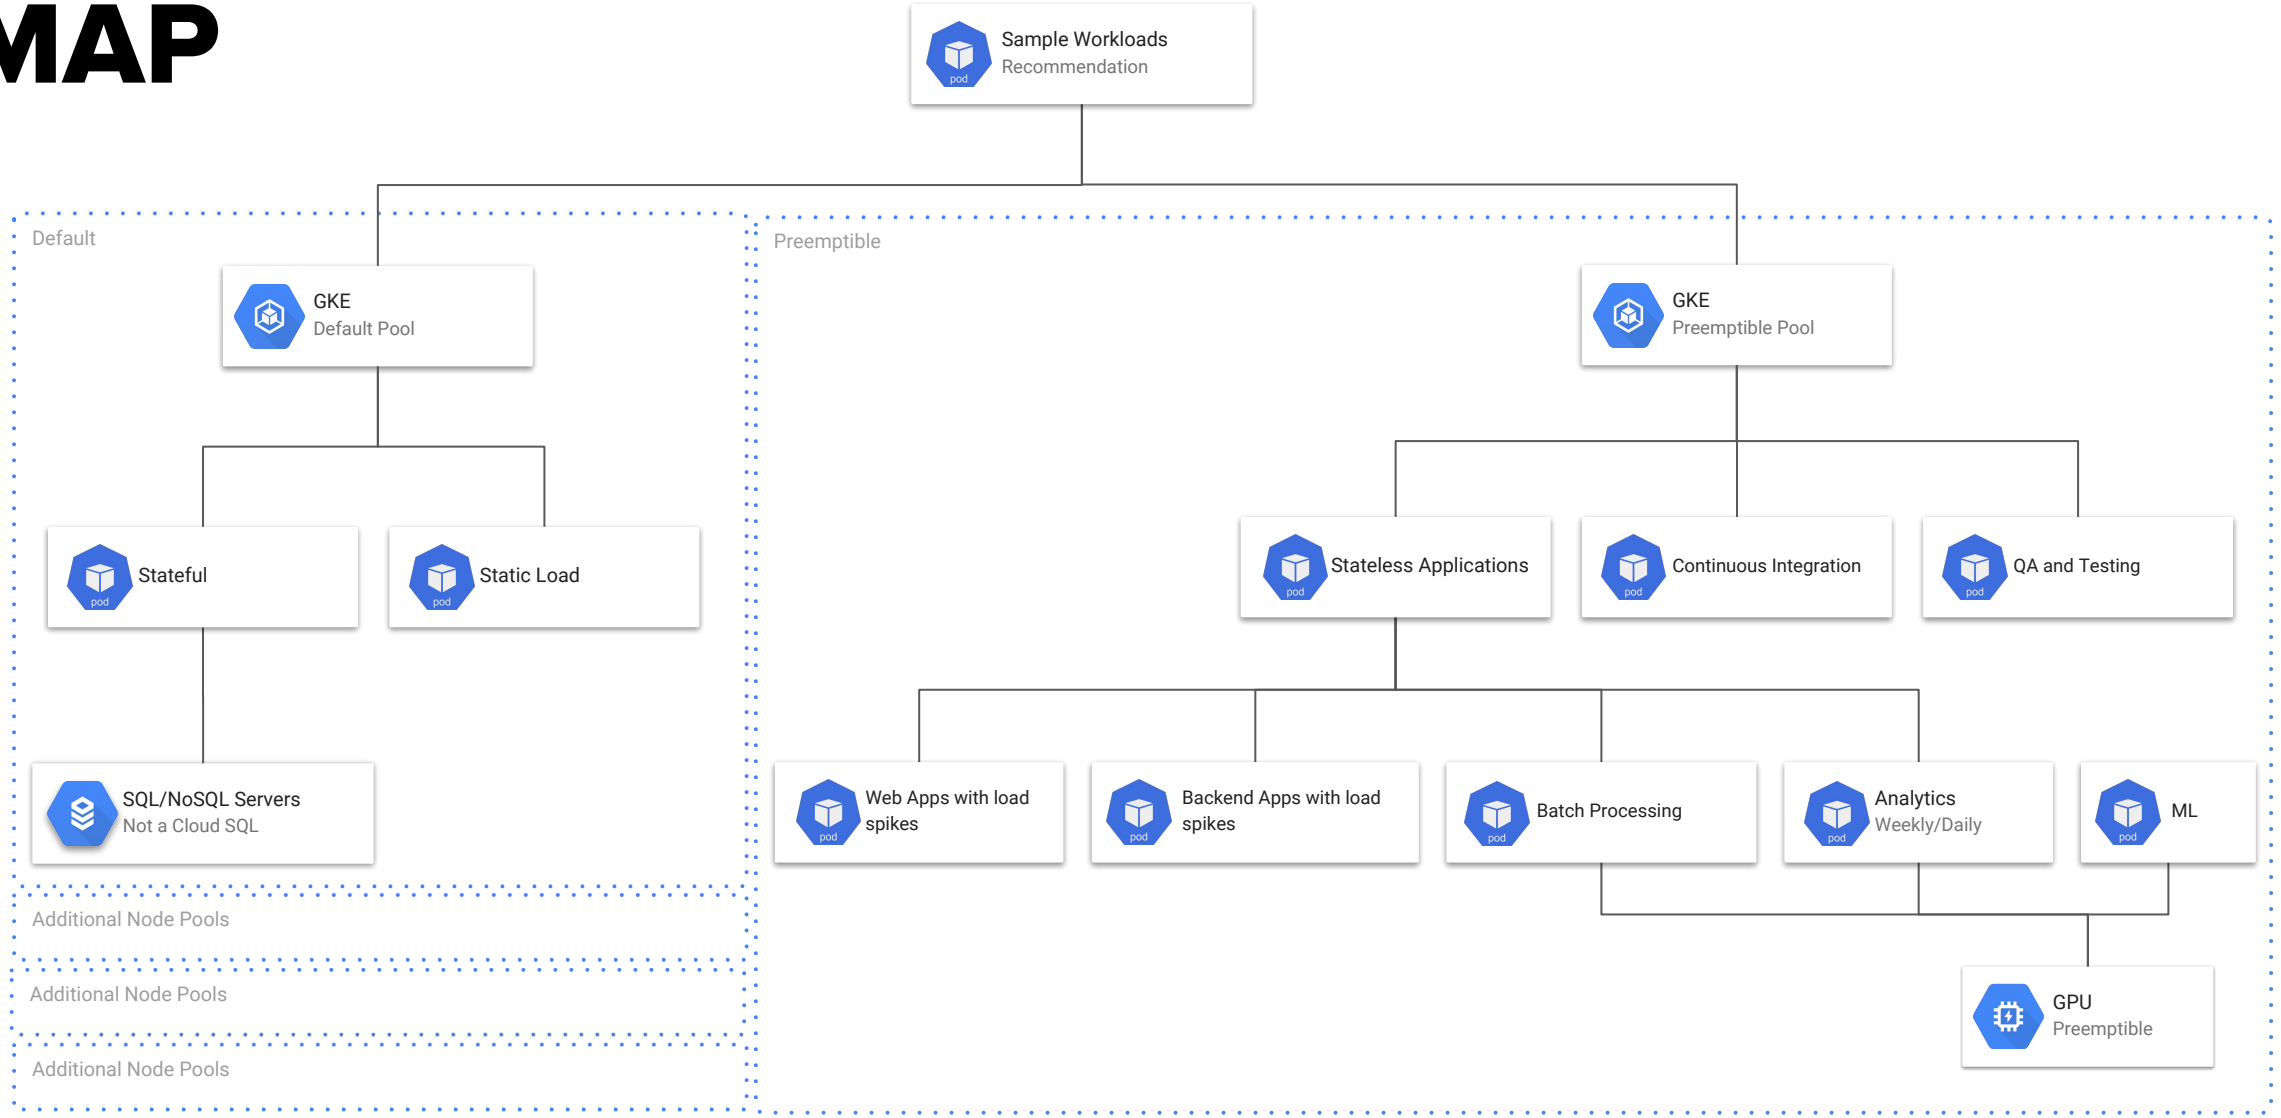

EXAMPLE 1

Simple App manager by Autoscaler

softserve

SAMPLE APP

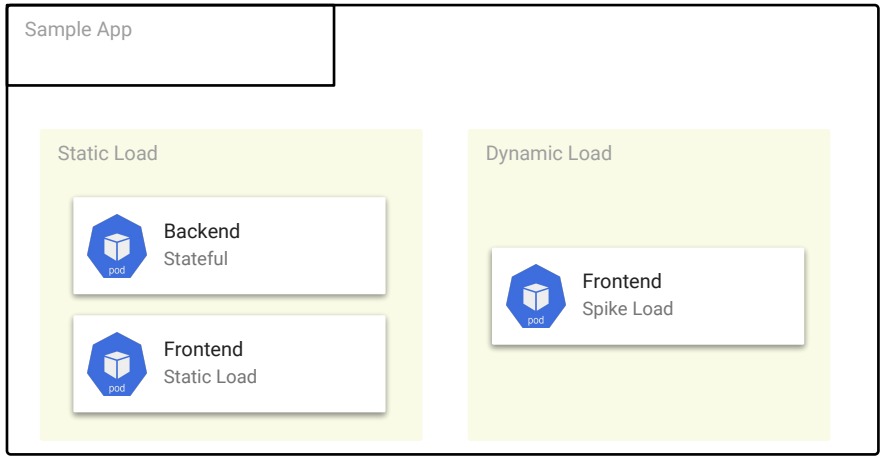

Old-style Java monolithic enterprise app splitted for two major parts

- **Backend Java** old-style stateful service
- New and shiny **Node.js Frontend** service
  - Predictable static load
  - High load spikes at EOD/EOW

Nodes in preemptible pool will have label

```
cloud.google.com/gke-preeptible: true
```

softserve

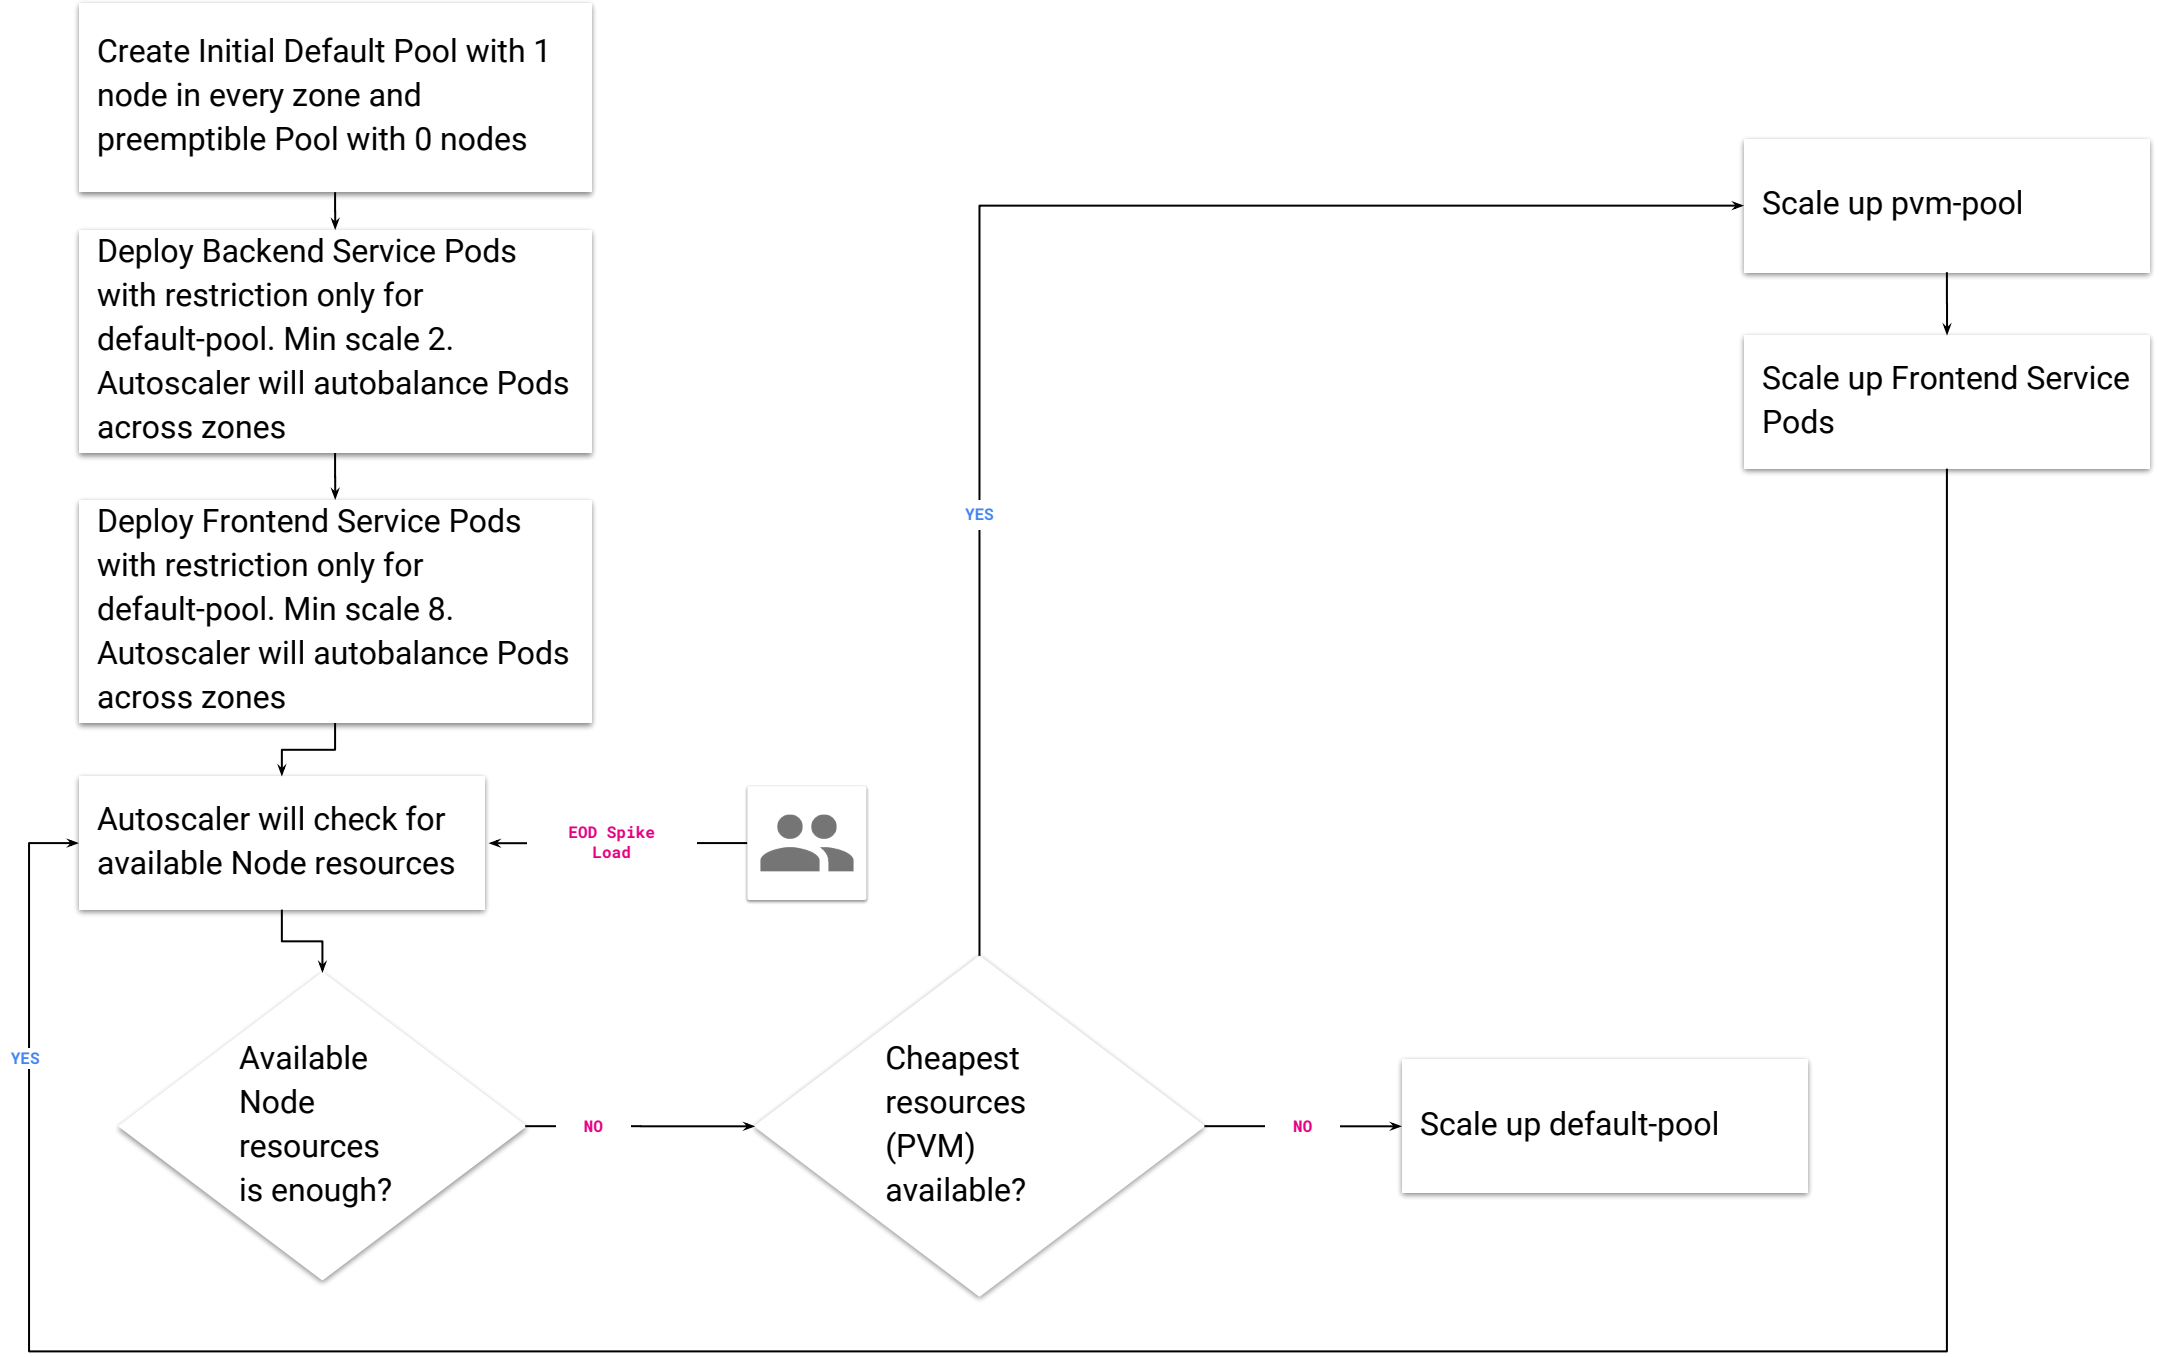
Create Initial Default Pool with 1 node in every zone and preemptible Pool with 0 nodes
Deploy Backend Service Pods with restriction only for default-pool. Min scale 2. Autoscaler will autobalance Pods across zones
Deploy Frontend Service Pods with restriction only for default-pool. Min scale 8. Autoscaler will autobalance Pods across zones
Autoscaler will check for available Node resources
EOD Spike Load
Available Node resources is enough?
YES
NO
Cheapest resources (PVM) available?
YES
NO
Scale up pvm-pool
Scale up Frontend Service Pods
Scale up default-pool

EXAMPLE 2

Let’s add Analytics

softserve

SAMPLE APP. ANALYTICS ADDED

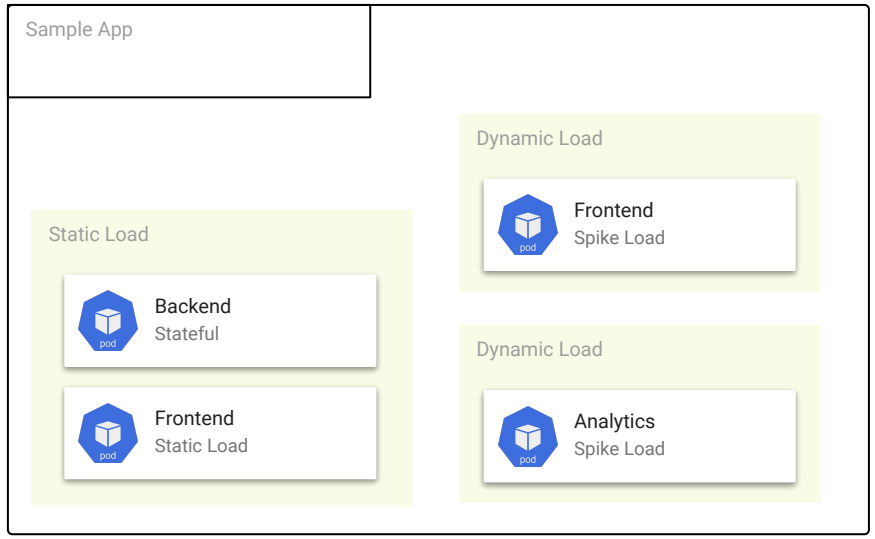

Old-style Java monolithic enterprise app splitted for two major parts

- Backend Java old-style stateful service
- New and shiny Node.js Frontend service
  - Predictable static load
  - High load spikes at EOD/EOW
- **Daily/Weekly Analytics**
  - To decrease cost let’s run heavy analytics only on PVM

softserve

GKE — MULTI-ZONE DEFAULT NODE POOL / MULTI-ZONE PREEMPTIBLE NODE POOL

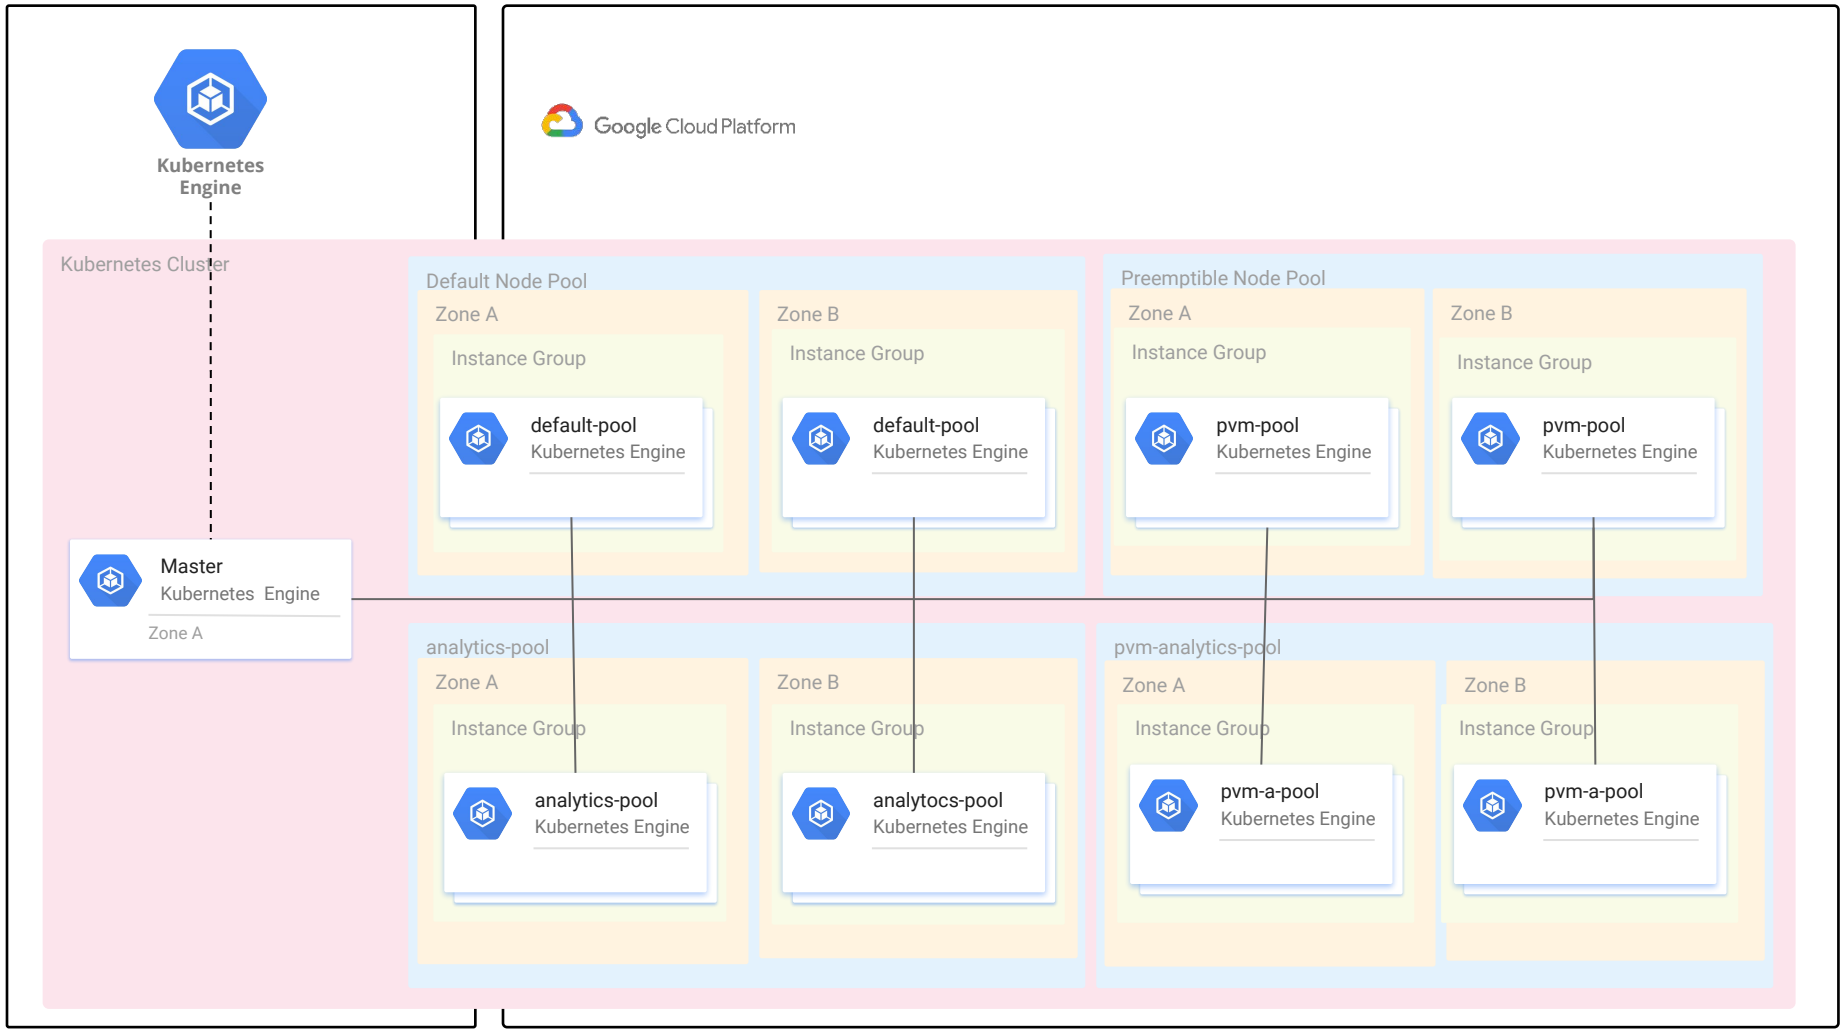

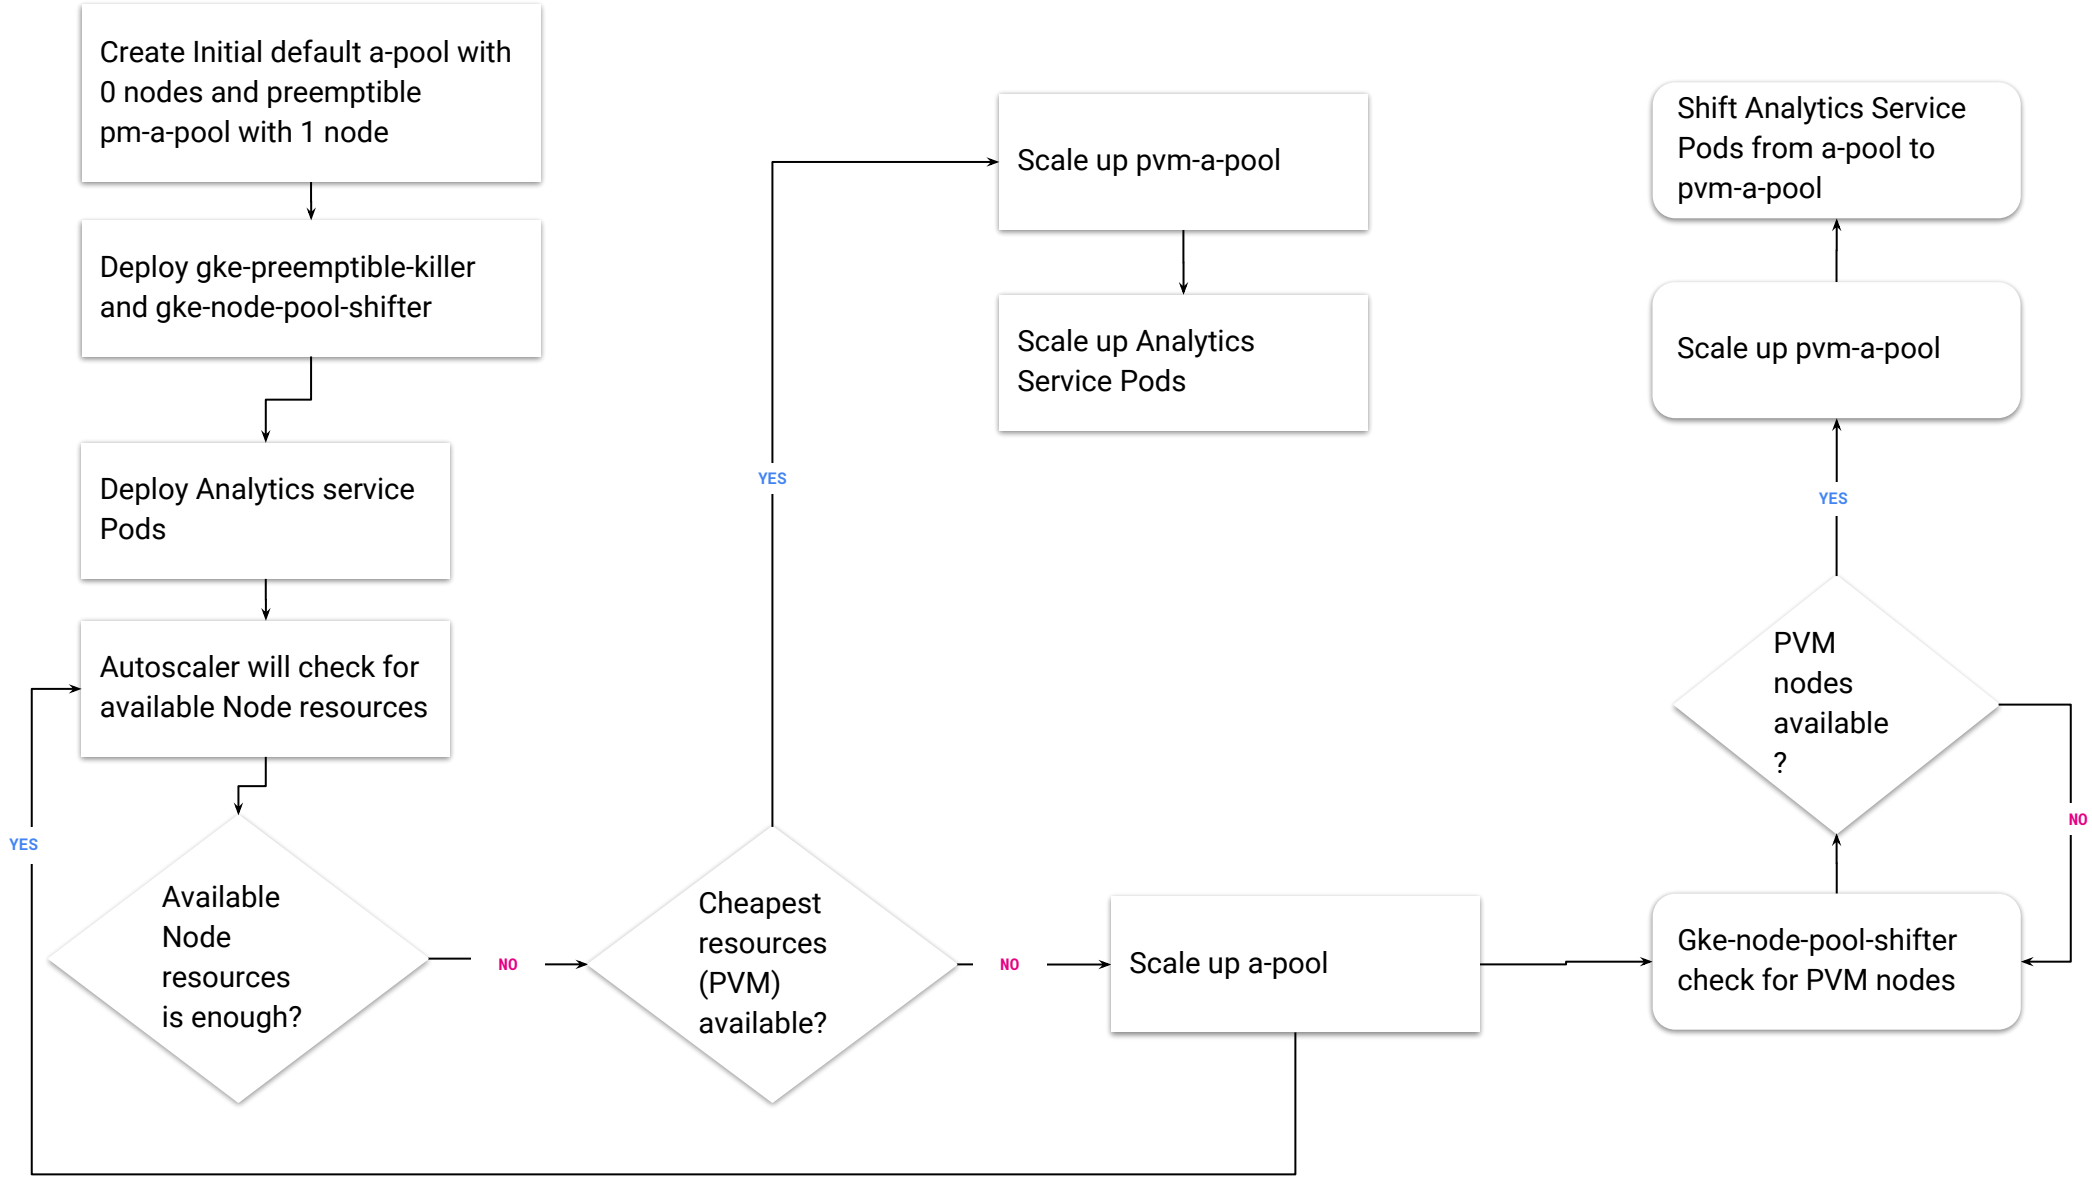
Create Initial default a-pool with 0 nodes and preemptible pm-a-pool with 1 node
Deploy gke-preemptible-killer and gke-node-pool-shifter
Deploy Analytics service Pods
Autoscaler will check for available Node resources
Available Node resources is enough?
YES
NO
Cheapest resources (PVM) available?
YES
Scale up pvm-a-pool
Scale up Analytics Service Pods
NO
Scale up a-pool
Gke-node-pool-shifter check for PVM nodes
PVM nodes available ?
YES
Scale up pvm-a-pool
Shift Analytics Service Pods from a-pool to pvm-a-pool
NO

SCHEDULE PODS ON PVM IF AVAILABLE

Modify your Pod or Deployment spec using

```
...
spec:
  affinity:
    nodeAffinity:
      preferredDuringSchedulingIgnoredDuringExecution:
      - preference:
          matchExpressions:
          - key: cloud.google.com/gke-preemptible
            operator: Exists
        weight: 100
...
```

SCHEDULE PODS ON NON-PVM NODES

Modify your Pod or Deployment spec using

```
...
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
        - matchExpressions:
          - key: cloud.google.com/gke-preemptible
            operator: DoesNotExist
...
```

CAPACITY VS COST(PVM)

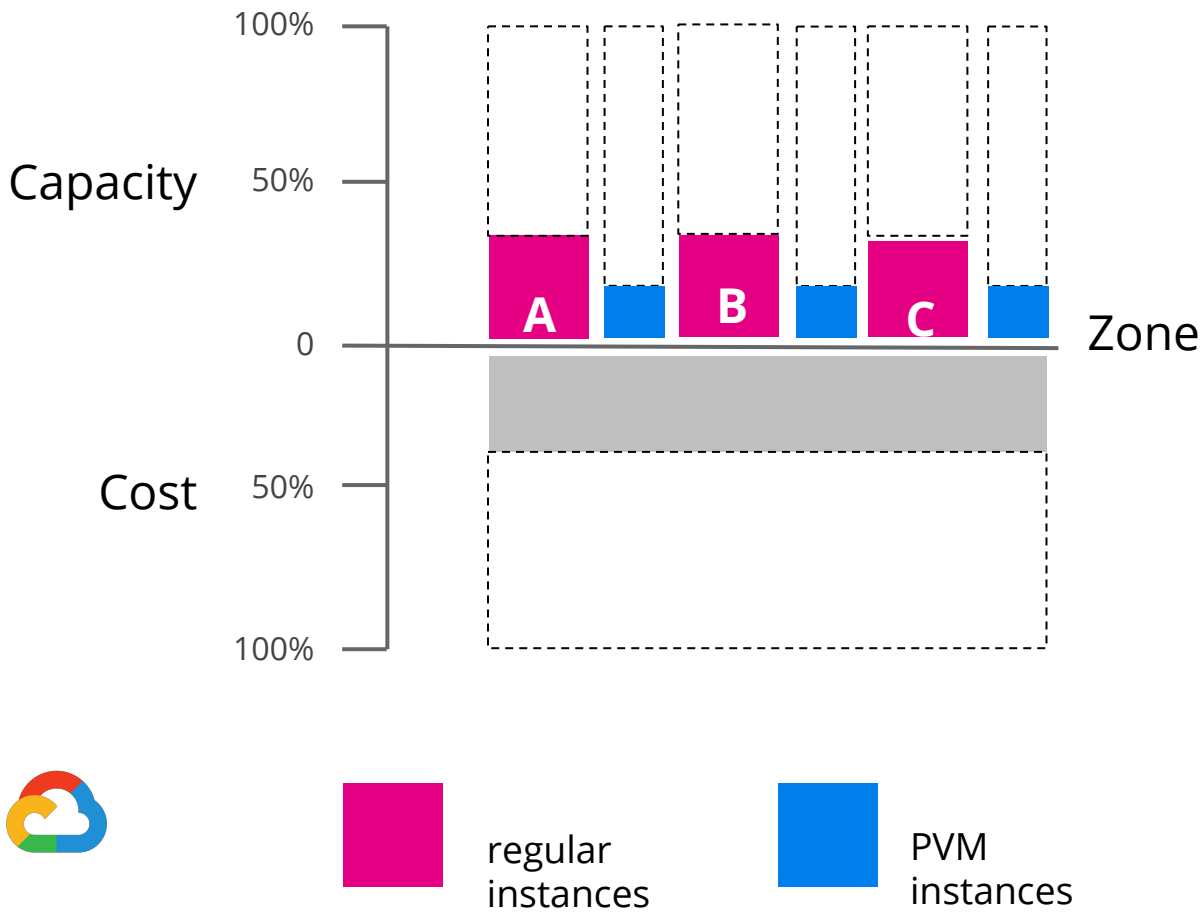

BEST CASE:

34% of spend,
100% of the time

WORST CASE:

120% of spend,
100% of the time
2x capacity

softserve

SHOW STACKDRIVER

softserve

DEMO STRESS

softserve

COST SAVING WITH PVM AS % OF CLUSTER

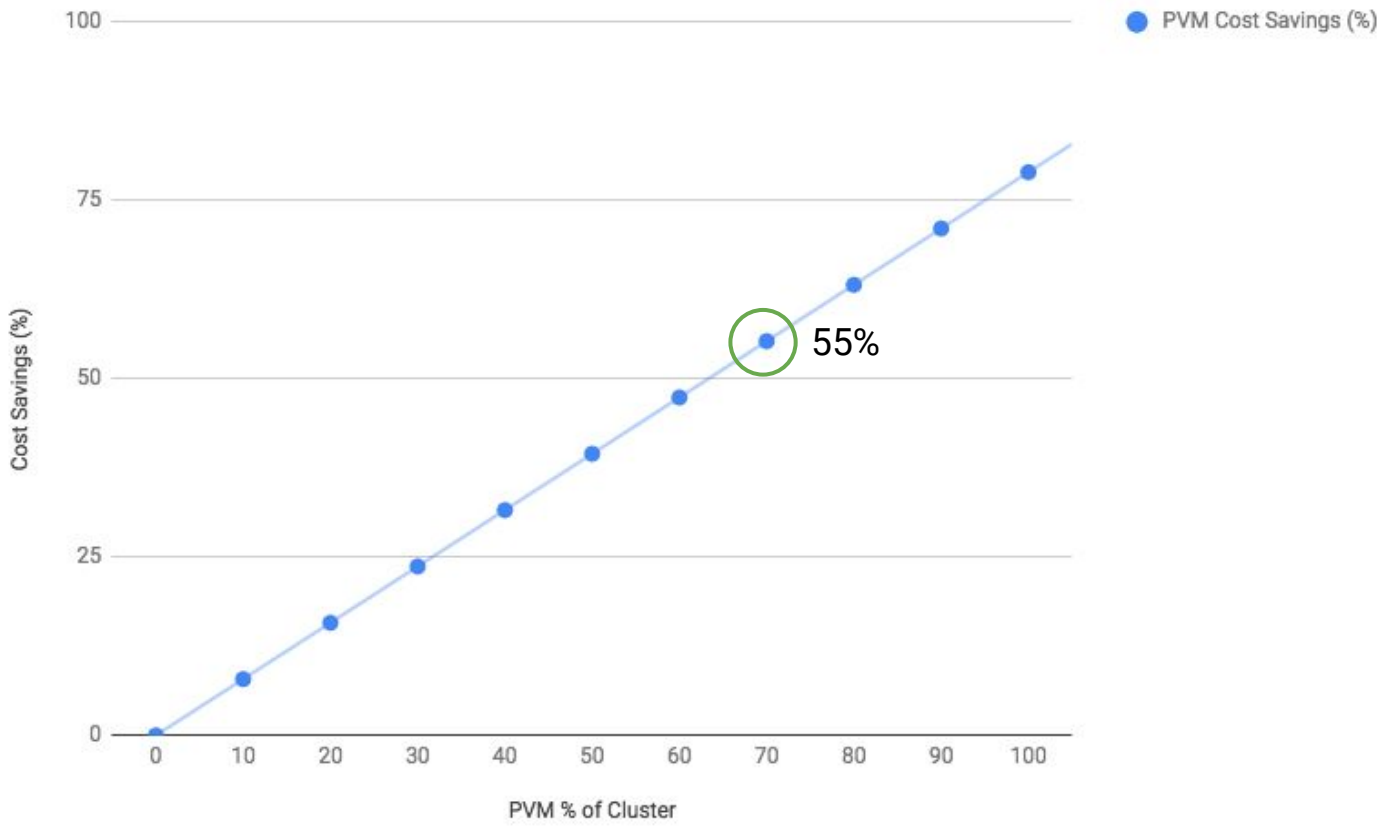

softserve

COMMITTED USE DISCOUNTS

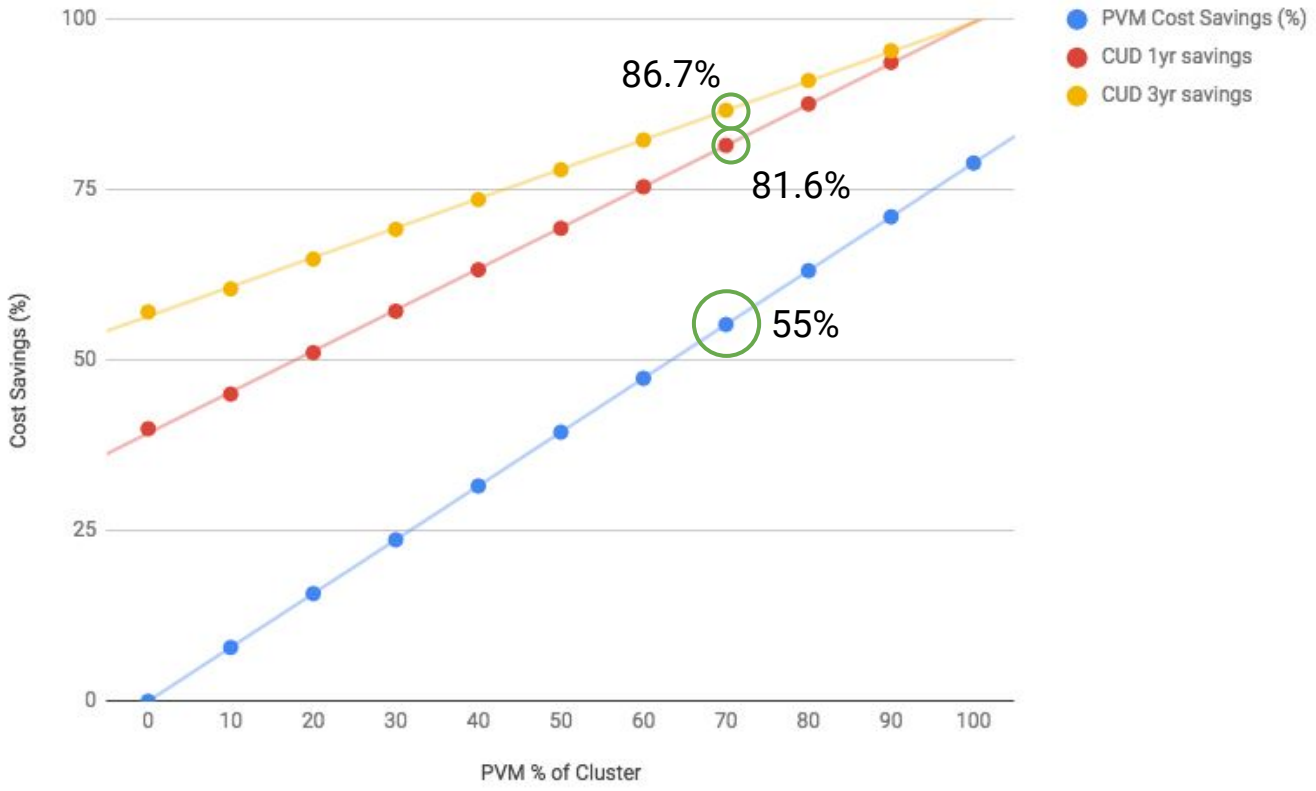

softserve

SHOW STACKDRIVER FOR STRESS

softserve

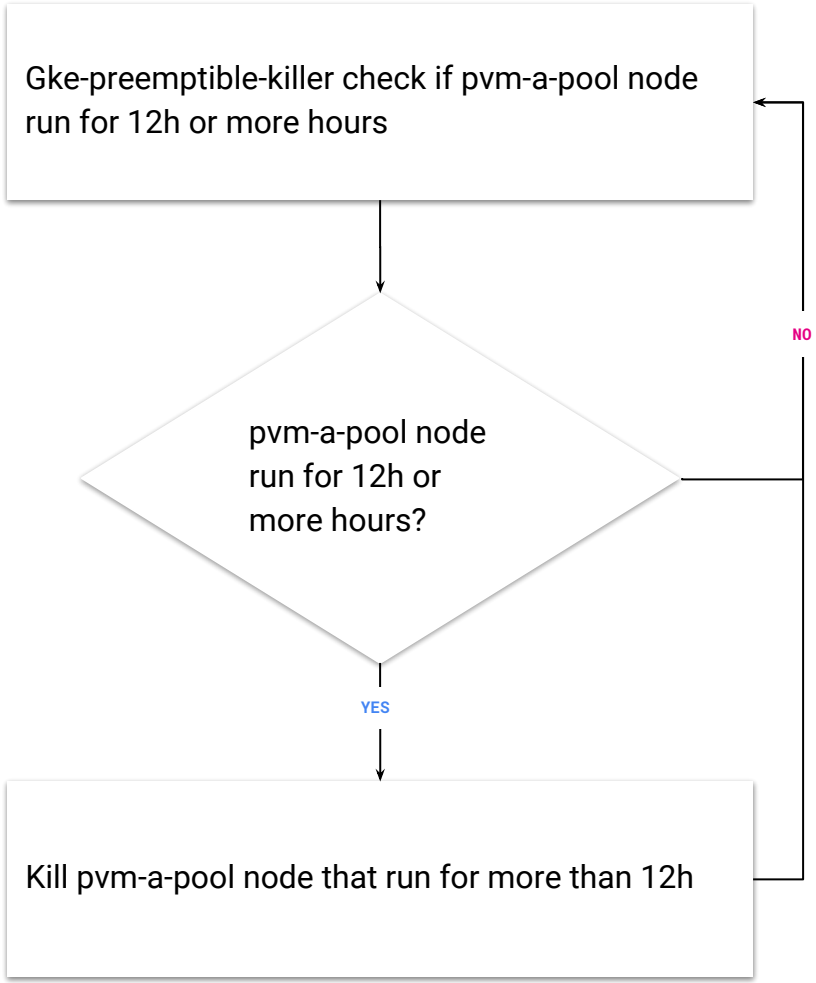

Kubernetes controller that ensures **deletion** of preemptible nodes in a GKE cluster **is spread out to avoid the risk** of all getting deleted at the same time after 24 hours

softserve

KUBERNETES CONTROLLERS FOR PREEMPTIBLE NODE POOL

- We recommend to randomly kill PVM in Preemptible Node Pool to avoid expire all nodes same time
  estafette-gke-preemptible-killer

- Another great tool will help to constantly monitor Node Pools and move nodes to preemptible PVM
  estafette-gke-node-pool-shifter

softserve


FOR
THE
FUTURE
softserve

APPENDIX

DEMO repository https://github.com/stonevil/gke-meetup-demo-project

softserve

KUBERNETES AND EPHEMERAL COMPUTE

Abstract

In cloud computing, elasticity is defined as "the degree to which a system is able to adapt to workload changes by provisioning and de-provisioning resources in an autonomic manner, such that at each point in time the available resources match the current demand as closely as possible".

For those familiar with Kubernetes, this may seem like a solved problem. Not exactly. What about the underlying cluster resources? Node Autoscaling is a great feature of GKE but how can we take advantage of this in interesting, cost effective way?

In this interactive session, you will walk through few cases how-to cut Google Cloud Kubernetes Engine cluster cost with preemptible VM’s, Stackdriver and Cluster Autoscaler.

softserve

BIOS

Ryan Richard — Ryan is a Customer Engineer at Google working with enterprise customers and focused on GCP. He has background building and running services in Kubernetes and originally added the Rackspace deployment code to the repo in 2014.

Myroslav Rys — Myroslav is a Solution Architect at SoftServe Inc. More than 8 years experience in large scale enterprise solutions including SaaS / Clouds solutions. Experience building products and solutions with Kubernetes from 2015.

softserve

UTILIZATION VS CAPACITY VS COST

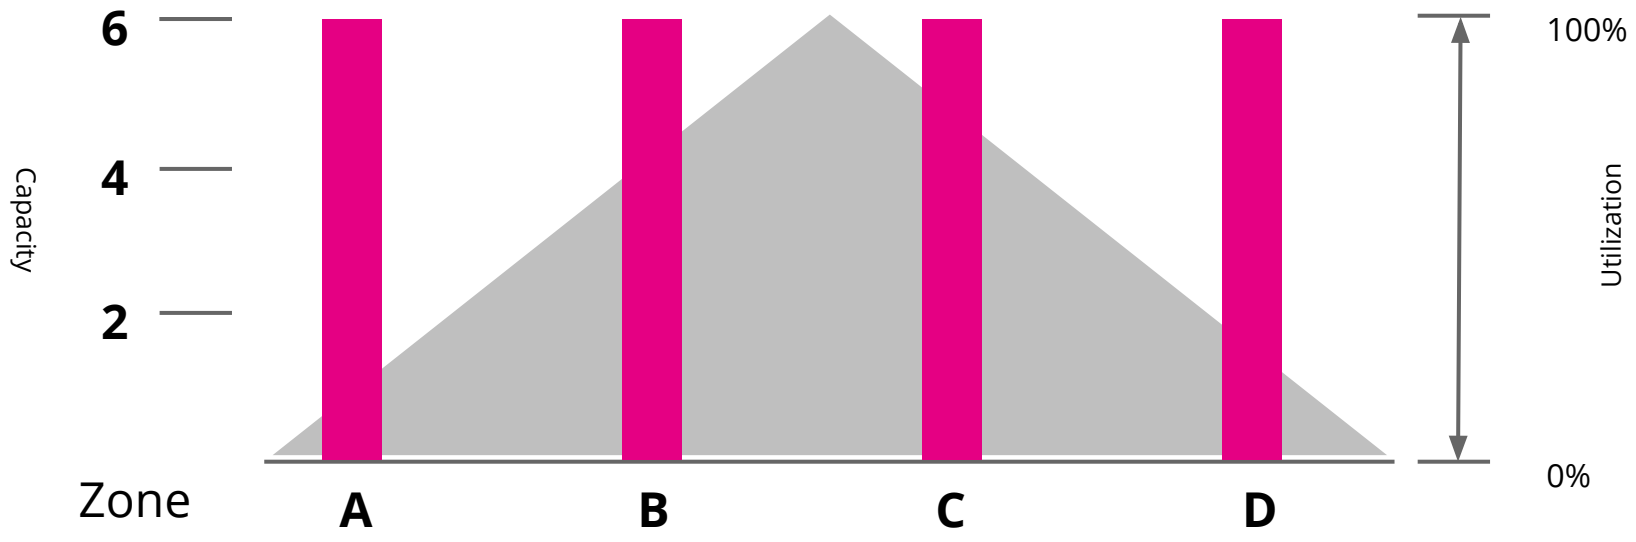

100% of spend,
100% of the time

softserve

SCALING (MYROSLAV)

Scaling

- Manually scaling, autoscaling (Cluster Autoscaler)
- Better utilization but you're paying full price for these resources
- **what if we look at this graph, it seems that our capacity spikes only for a few hours a day (Stackdriver)**
- Or a known batch job that will increase usage for a known amount of time, want to pay the least for it.
- Is there a way to handle this capacity temporarily without paying full price?

Cluster autoscaler considers the relative cost of each instance type in the node pool and attempts to expand the least expensive possible node pool. [1]

Cluster Autoscaler only takes requested resources into account for autoscaling. It does not take current utilization into account.

[1] https://cloud.google.com/kubernetes-engine/docs/concepts/cluster-autoscaler

softserve

UTILIZATION VS CAPACITY VS COST

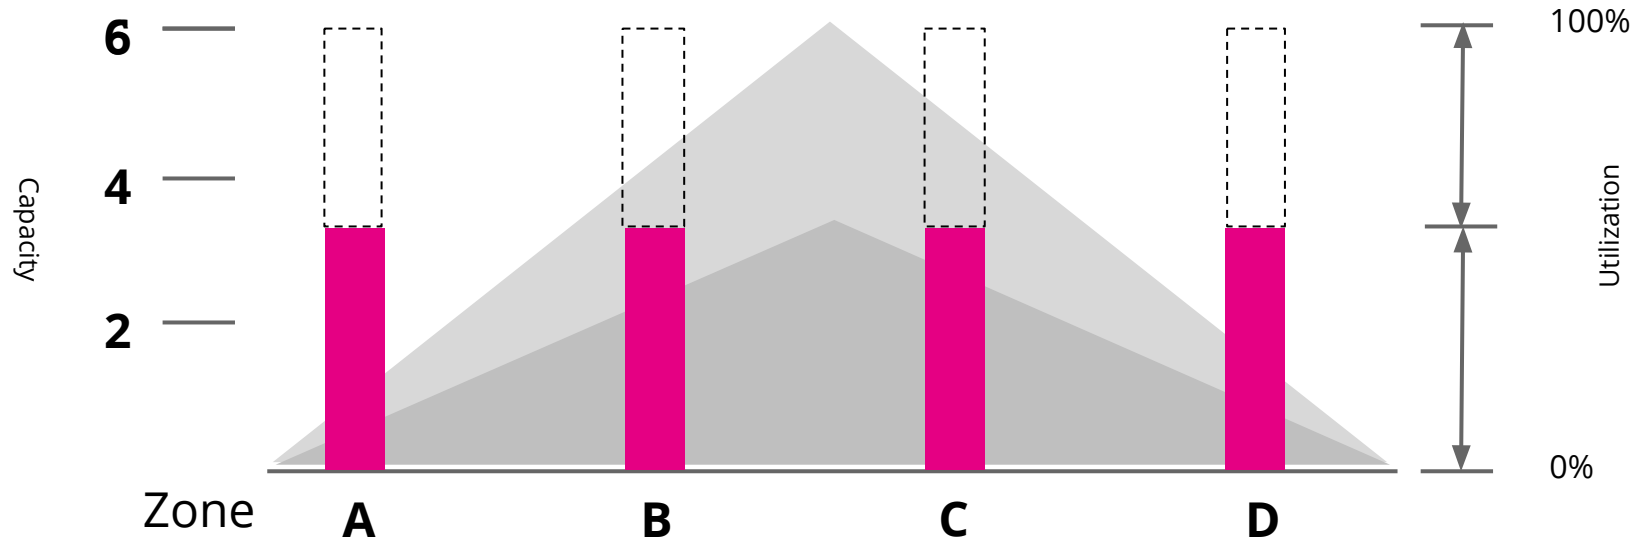

BEST CASE:

50% of spend,
100% of the time

WORST CASE:

100% of spend,
100% of the time

softserve

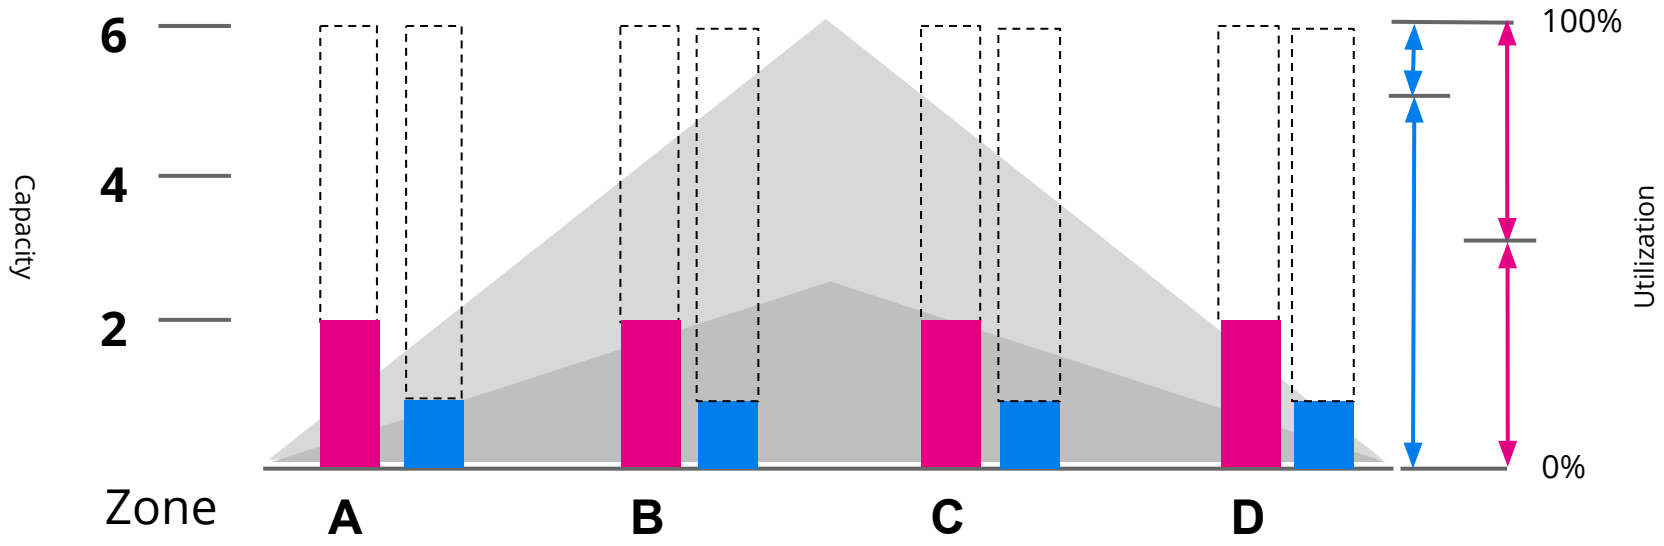

BEST CASE:

36.5% of spend,
100% of the time

WORST CASE:

120% of spend,
100% of the time
2x capacity

softserve

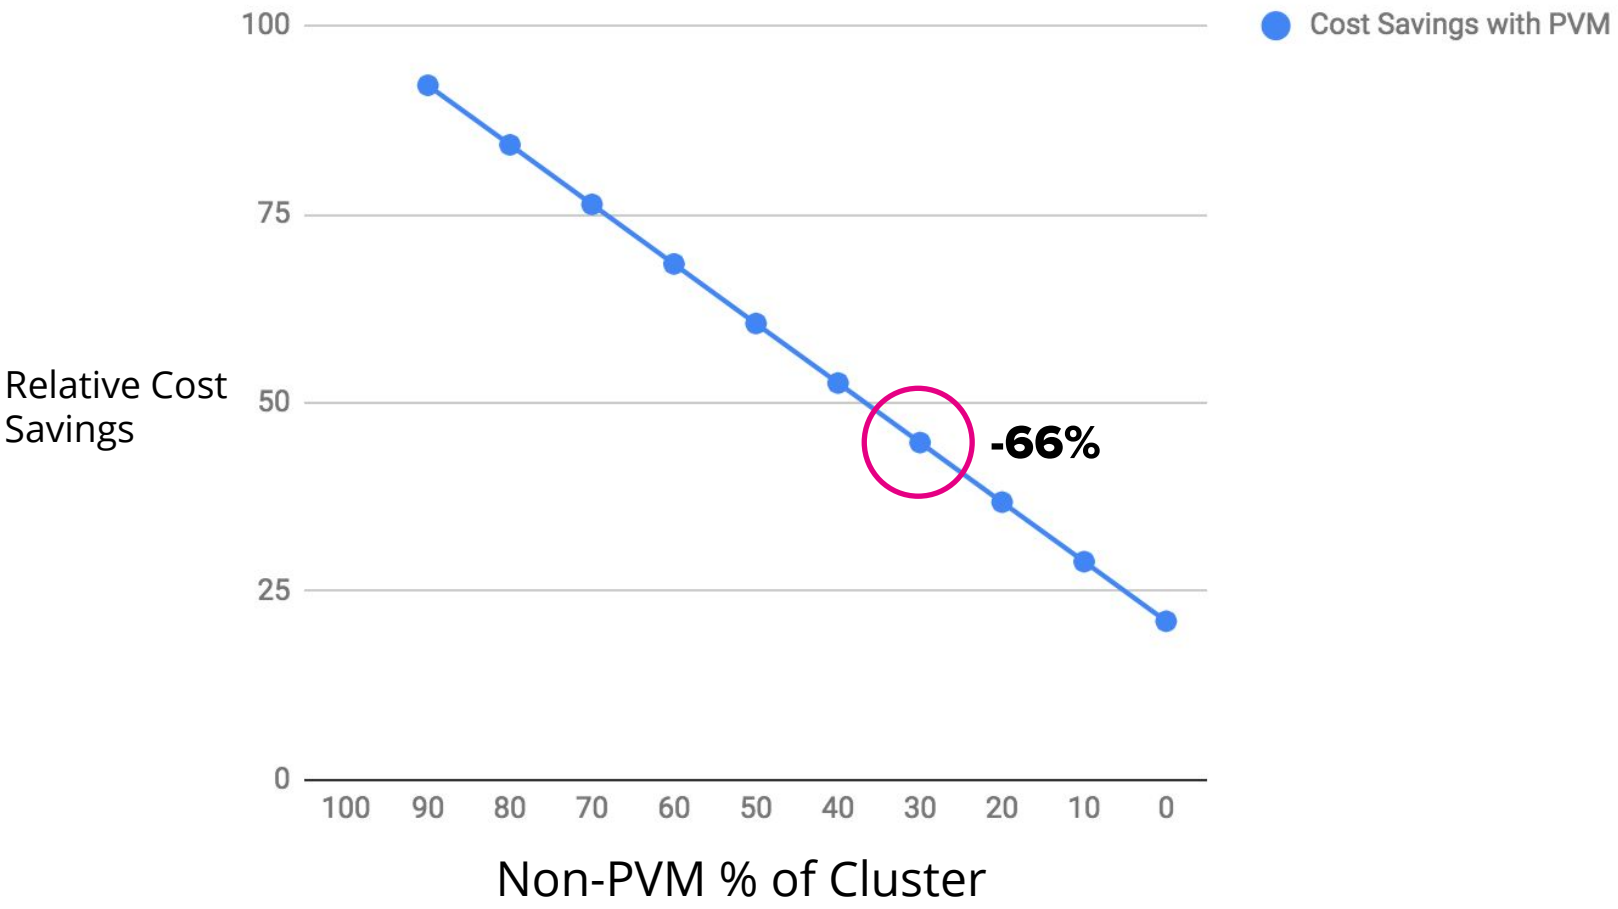
Cost Savings with PVM
100
75
50
25
0
Relative Cost Savings
-66%
100
90
80
70
60
50
40
30
20
10
0
Non-PVM % of Cluster

softserve

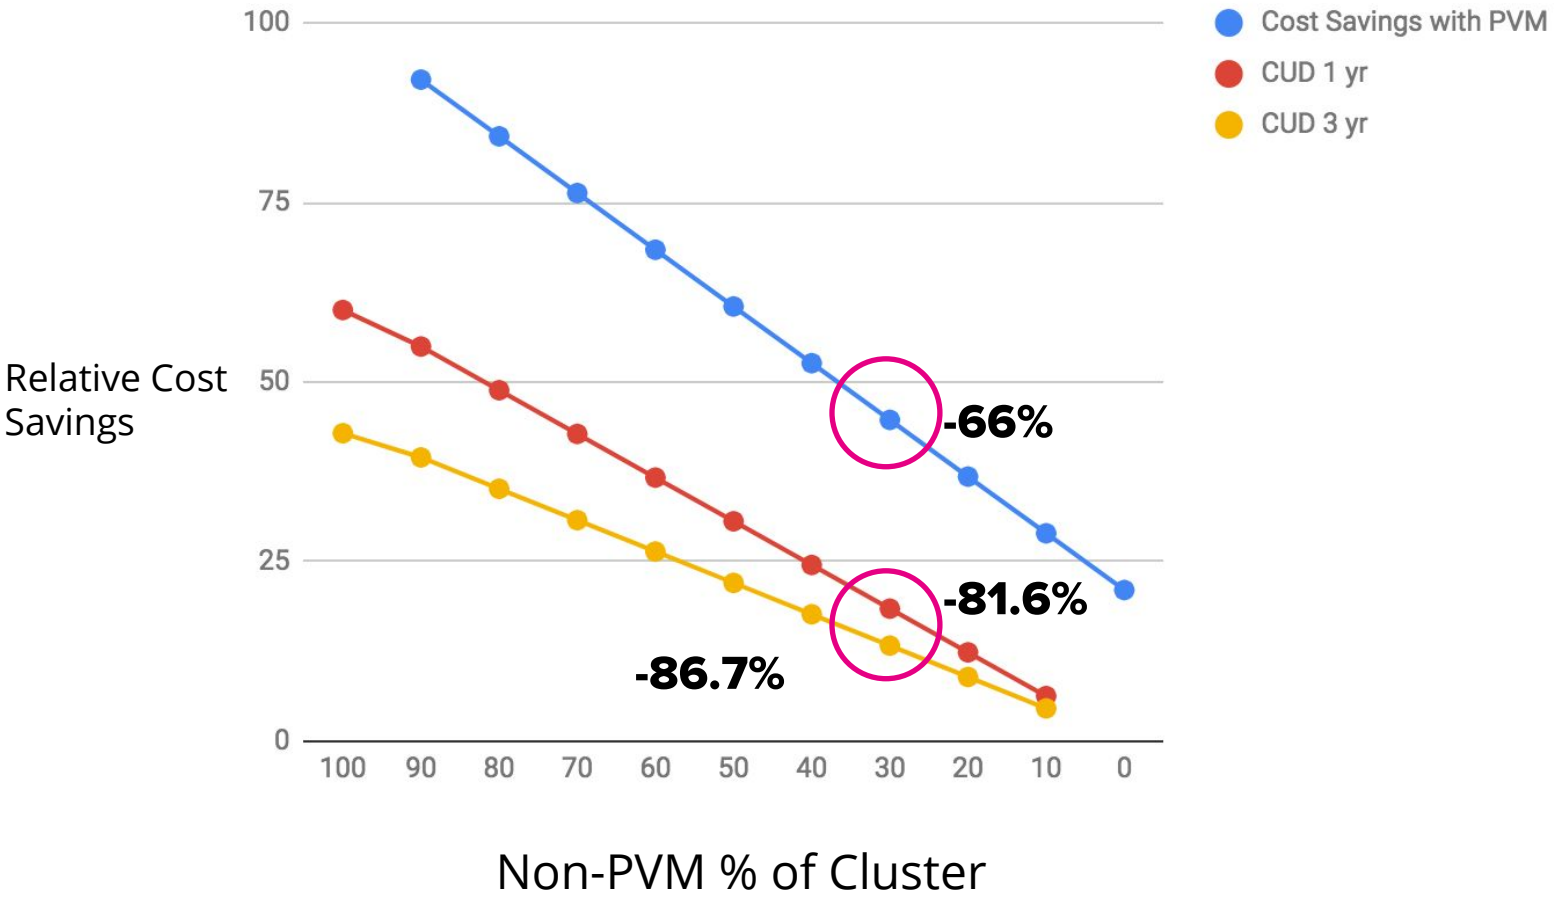
Relative Cost Savings
100
75
50
25
0
100 90 80 70 60 50 40 30 20 10 0
Cost Savings with PVM
CUD 1 yr
CUD 3 yr
-66%
-81.6%
-86.7%
Non-PVM % of Cluster

softserve